

Advanced Machine Learning

Lecture 2: Supervised Deep Learning

Syllabus

- Deep Sequential Modeling
 - RNN
 - LSTM
 - GRU
 - Bidirectional RNNs
- Convolutional NN (CNN)
 - What are they for?
 - Step 1: Convolution
 - Step 1(b) - ReLU
 - Step 2 - Pooling
 - Step 3 - Flattening
 - Step 4 - Full Connection

Deep Sequence modelling

Why do we need sequence modelling?

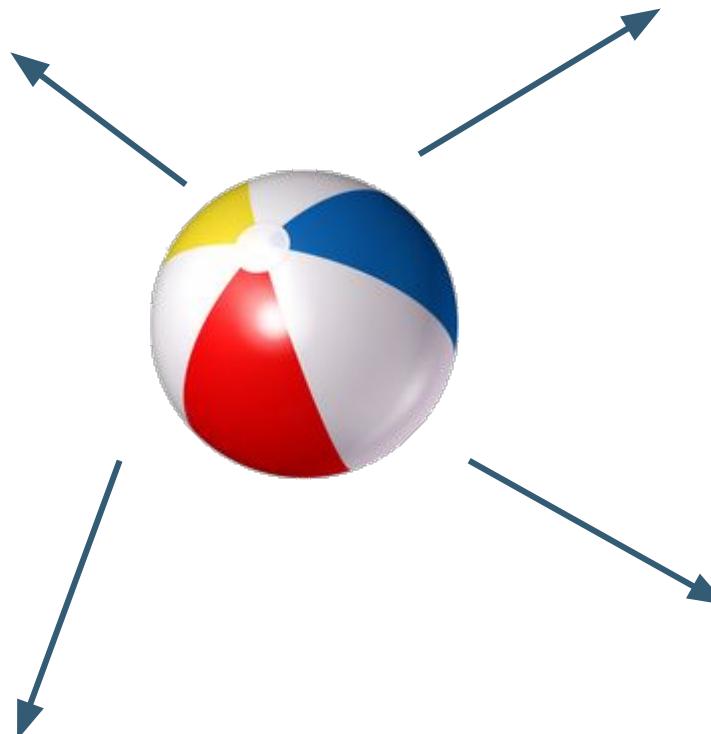
I Can assure, this ball is moving:
Can you predict the direction this ball is heading?



Why do we need sequence modelling?

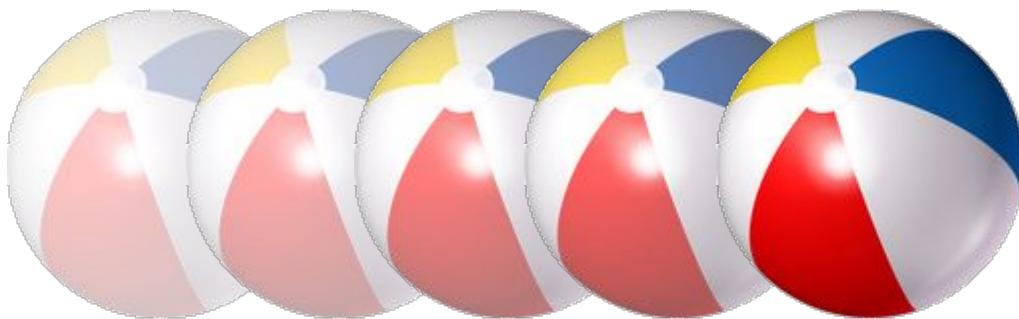
Can you predict the direction this ball is heading?

With no other information, it can be ANYWHERE



Why do we need sequence modelling?

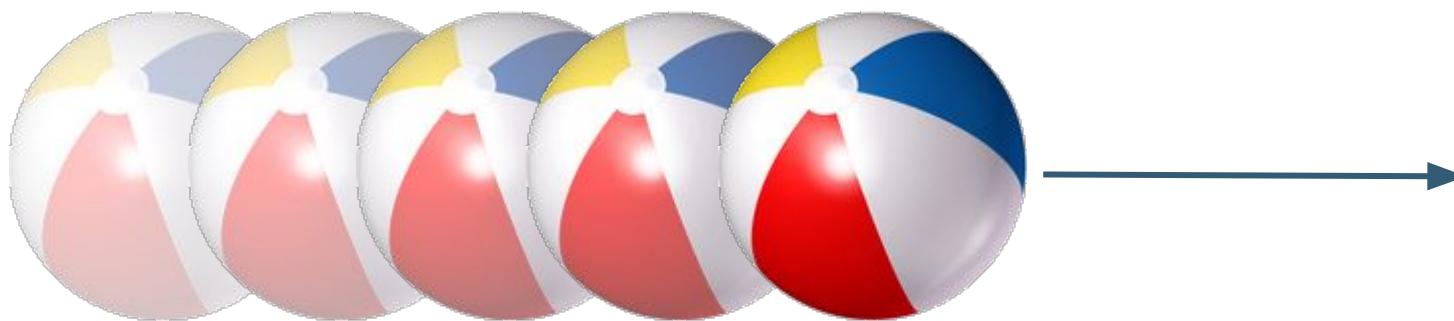
Can you predict the direction this ball is heading, if I tell you that its last movements were these ones?



Why do we need sequence modelling?

Can you predict the direction this ball is heading, if I tell you that its last movements were these ones?

It is easy.... it is heading East



Why do we need sequence modelling?

- Basically, in order to predict, sometimes, **past knowledge is needed.**
- Examples of Sequence data (or TimeSeries):
 - Audio or sensor data, is a sequence of values
 - Text can be seen as sequences of characters or words



Audio, Sensor data, ...

character sequence:
H E L L O M Y N A M E I S E K H I

Word sequence:
Hello My Name Is Ekhi

Text data

Can we predict Sequences using FFNN?

“This morning I took my cat for a walk.”

given these words

predict the
next word

- Limitation:
 - Inputs in FFNN are fixed, so we would not be able to predict variable length sentences.

“This morning I took my cat for a walk.”

given these two words predict the next word

- OK, let us take only 2 words and try to predict the 3rd one
- Limitation:
 - We loose long time dependency:

I am **Basque**, right now I live in Guatemala, but I still speak fluent _____



The food was good, not bad at all.

vs.

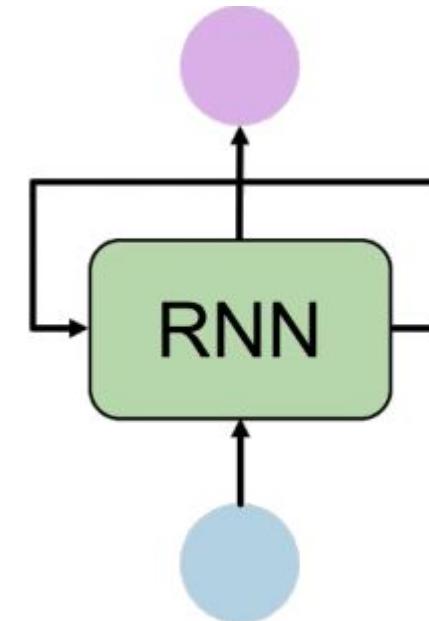


The food was bad, not good at all.

- There are even more Limitations:
 - Order importance
 - No information Share across the sequence

Sequence Modelling: Design Criteria

1. Handle Variable length Sequences
2. Track long-term dependencies
3. Maintain order information
4. Share parameters across the sequence



RNN in action - SunsPring



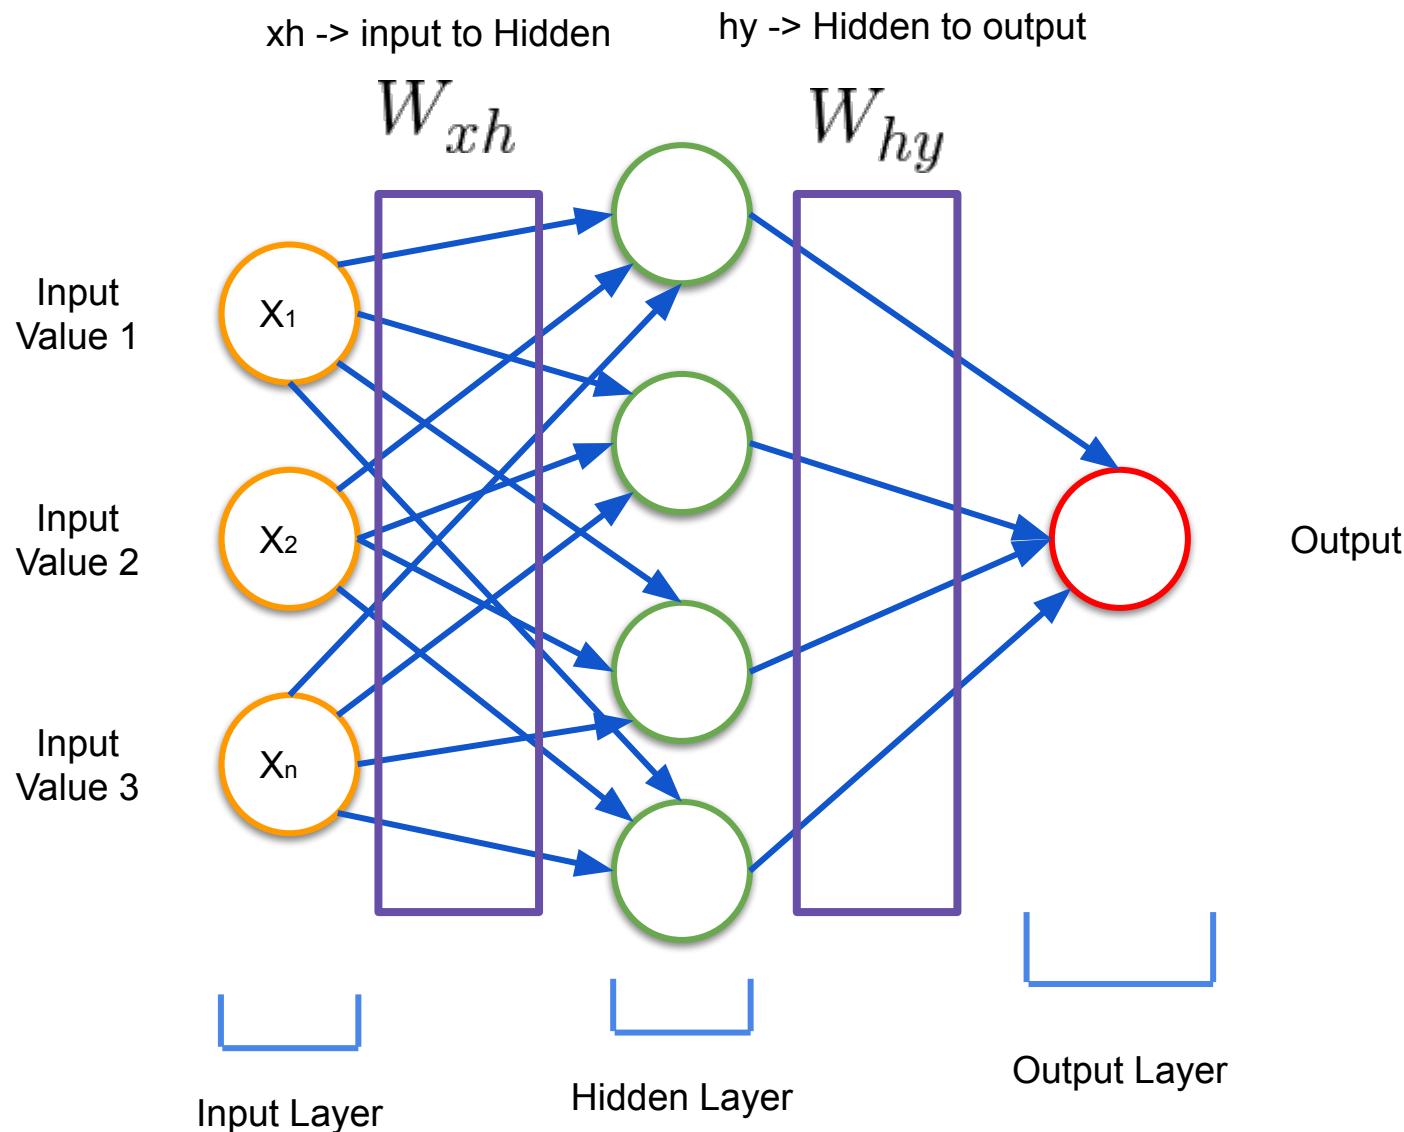
<https://www.imdb.com/title/tt5794766/>

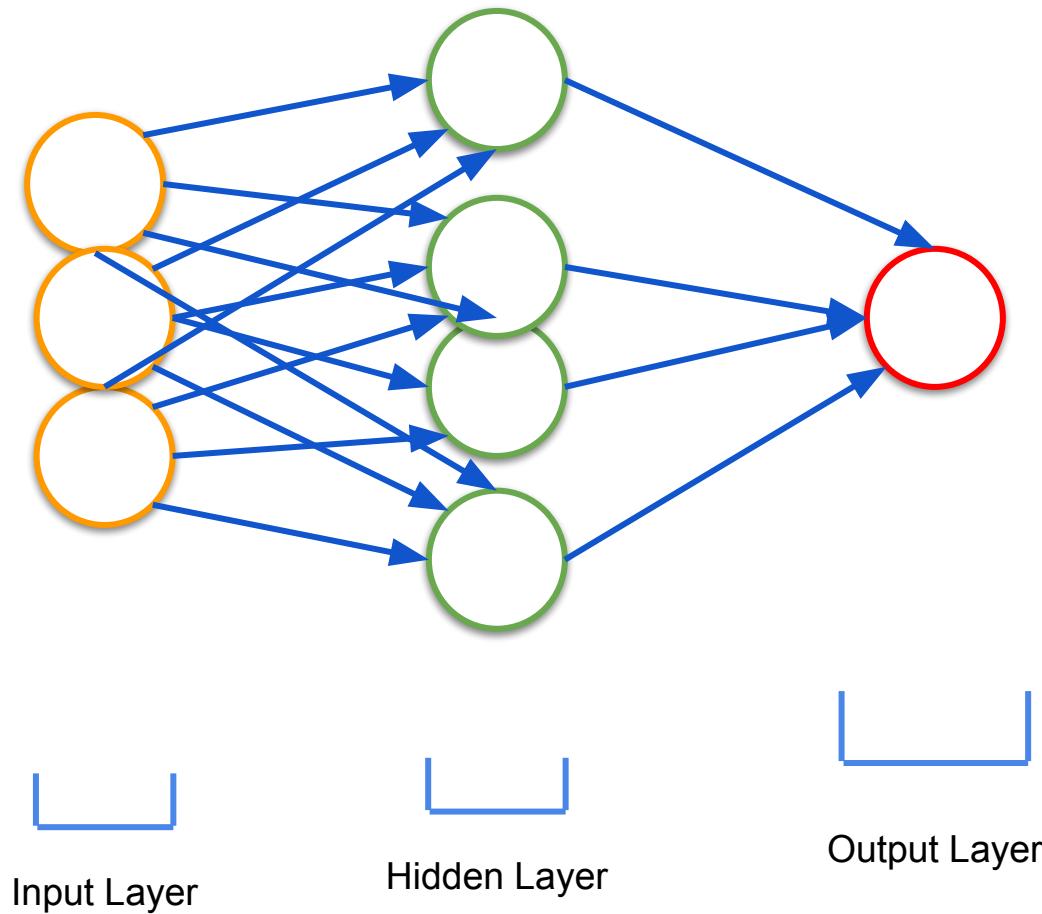
Short movie written by an RNN (llamado Benjamin)

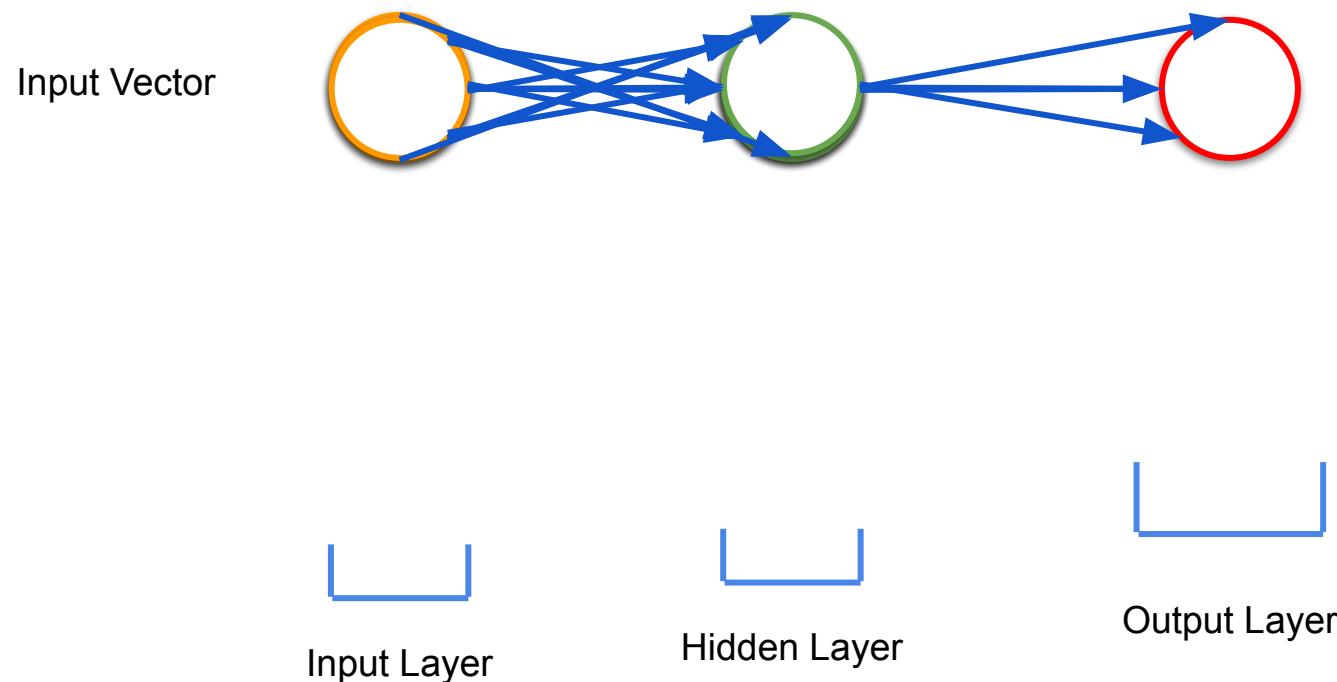
Learning: Hundreds of science fiction film plots
although each phrase has a meaning, there is a lack of a global vision

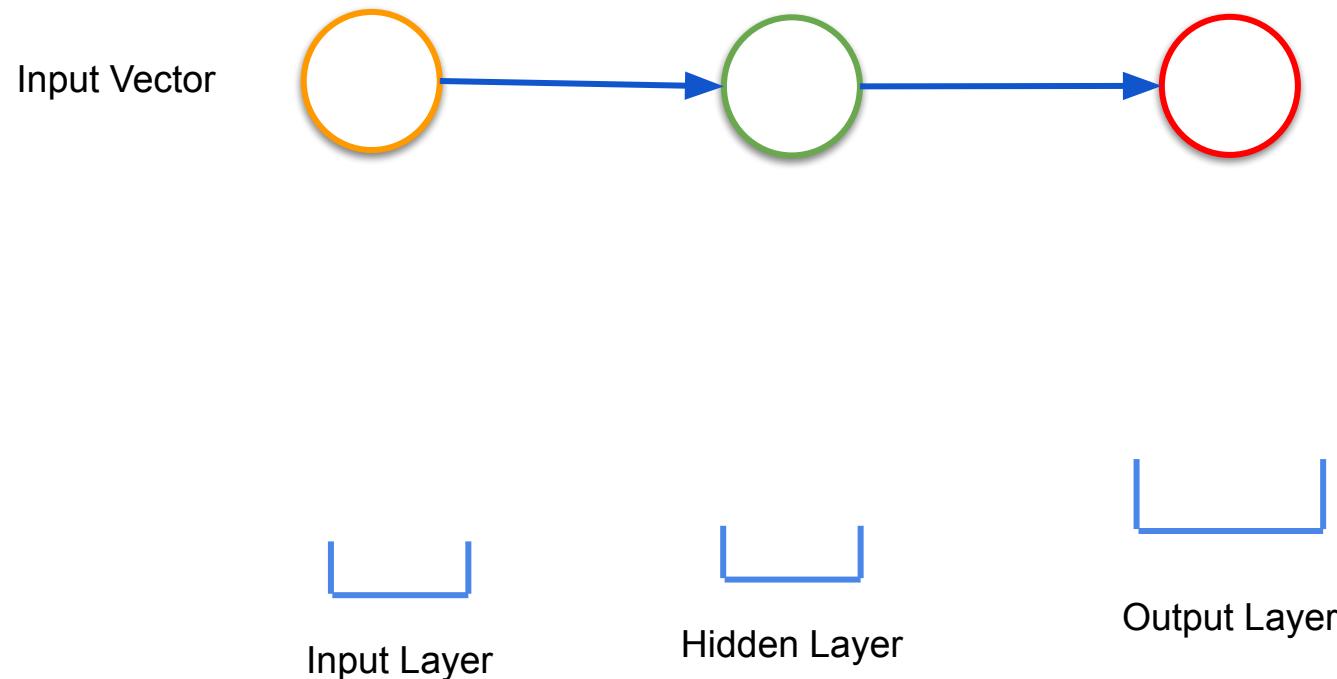
Recurrent Neural Networks

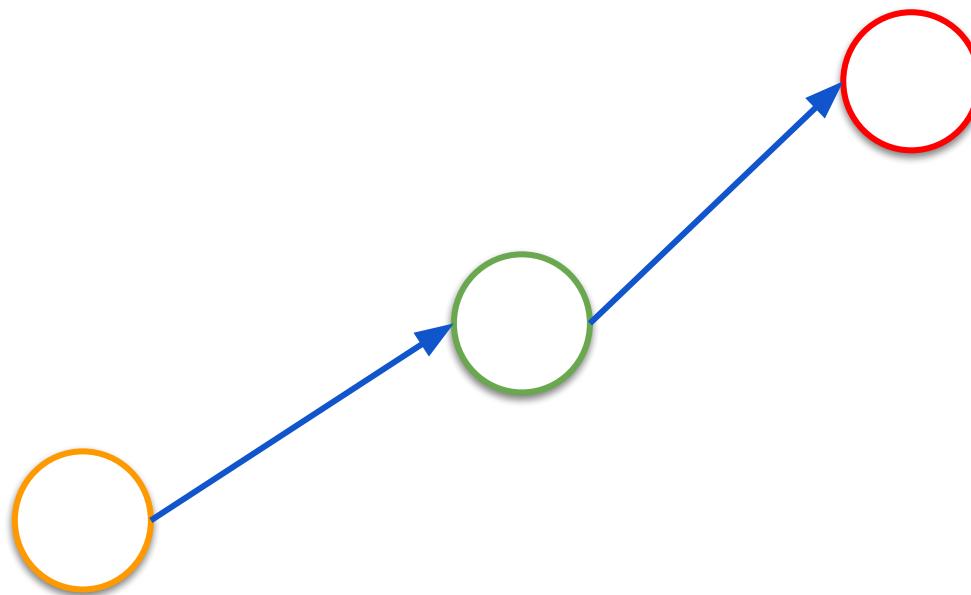
Recurrent Neural Networks (RNN)

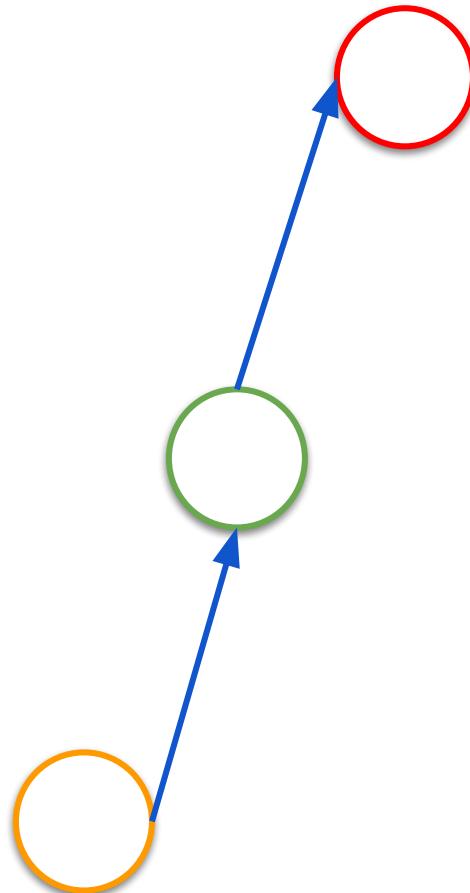


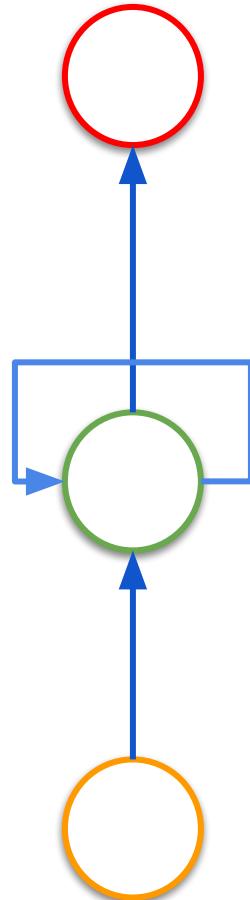




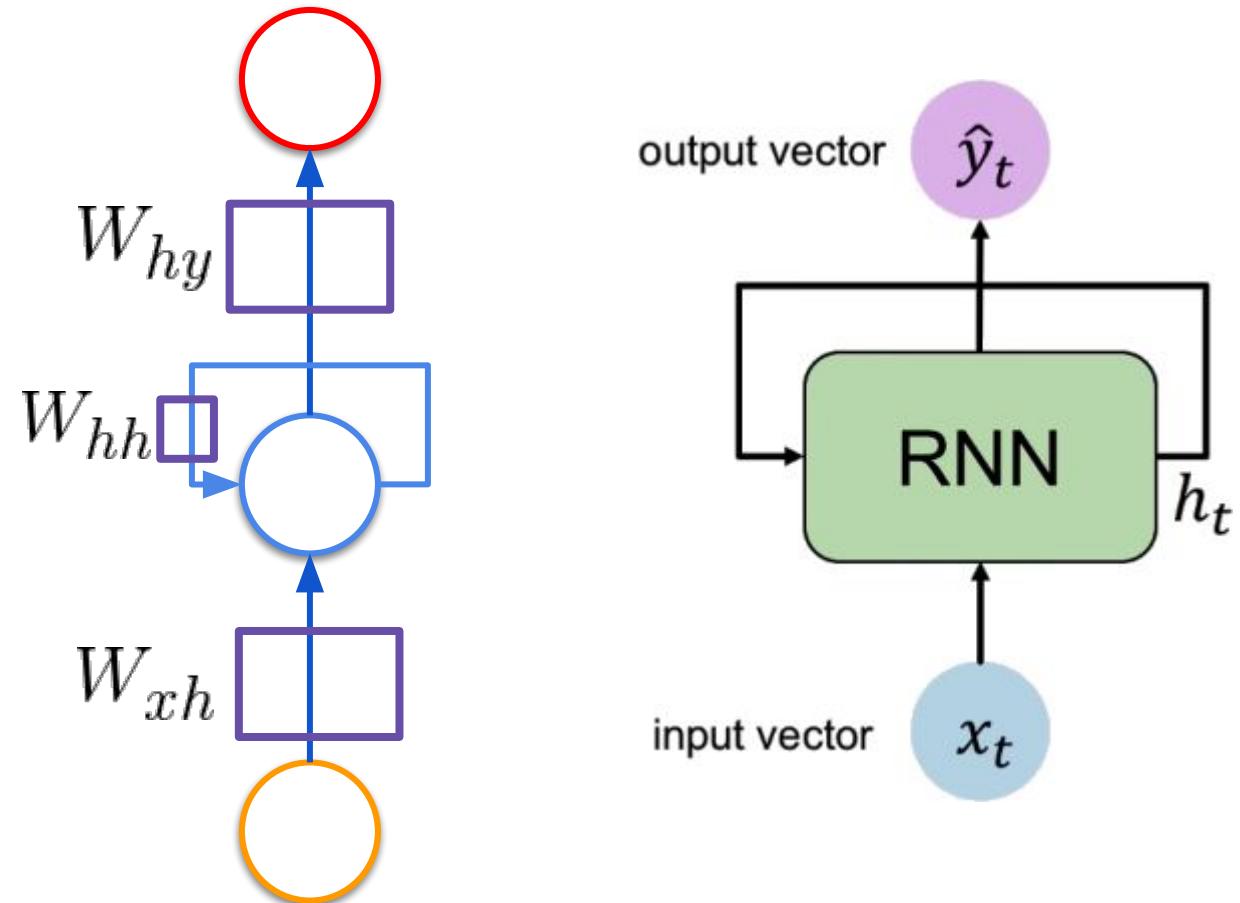




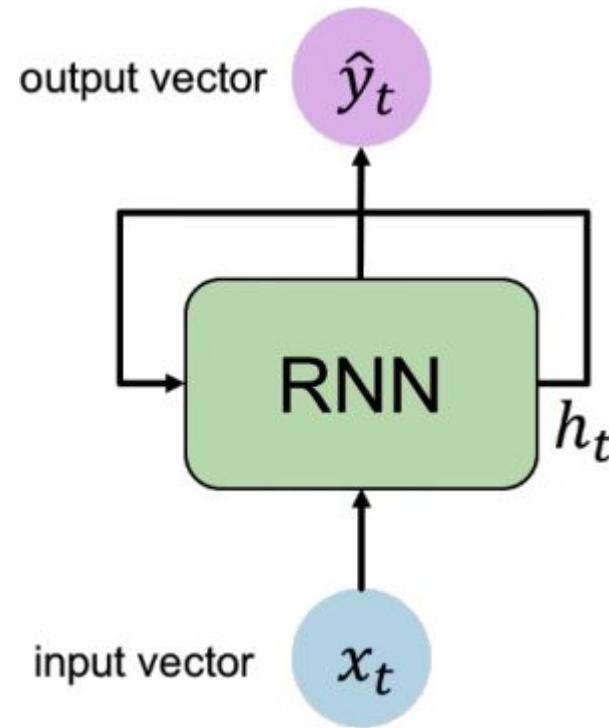
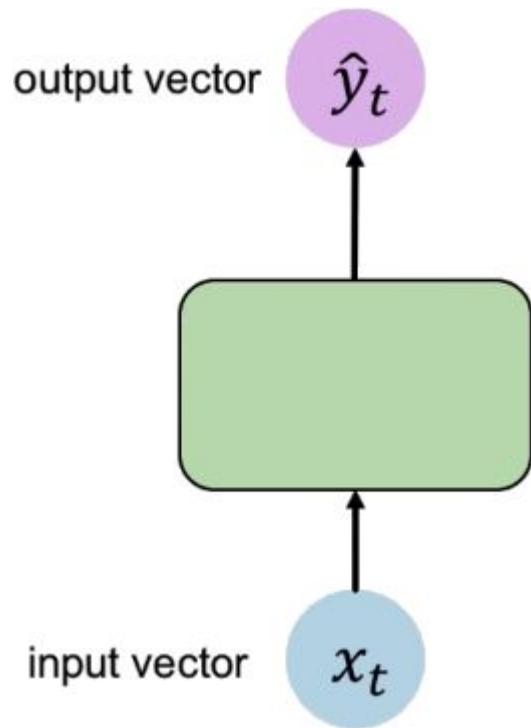




RNN



“Vanilla” FFNN vs RNN



- In each timeStamp, RNNs have as input:
 - x_t (current reading)
 - h_t (internal state)
- and one output:
 - \hat{y}_t

Output Vector

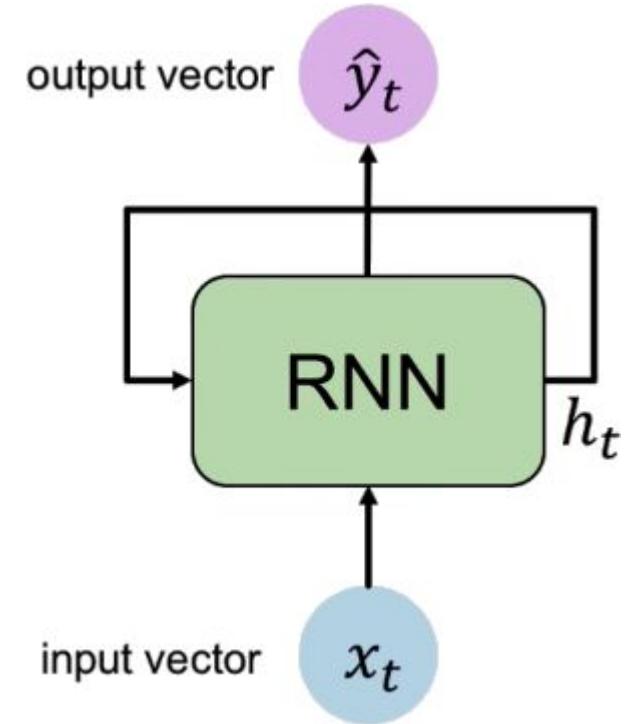
$$\hat{y}_t = \mathbf{W}_{hy}^T h_t$$

Update Hidden State

$$h_t = \tanh(\mathbf{W}_{hh}^T h_{t-1} + \mathbf{W}_{xh}^T x_t)$$

Input Vector

$$x_t$$



W_{xh} -> input to hidden
 W_{hh} -> hidden to hidden
 W_{hy} -> hidden to output

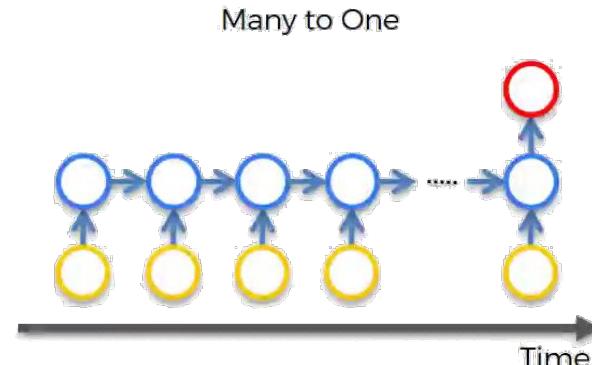
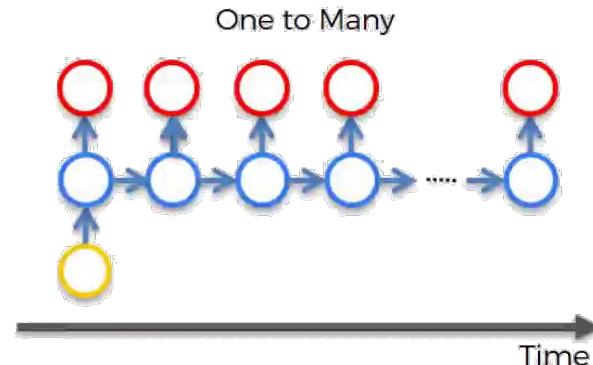
```
tf.keras.layers.SimpleRNN(rnn_units)
```



RNN Usages

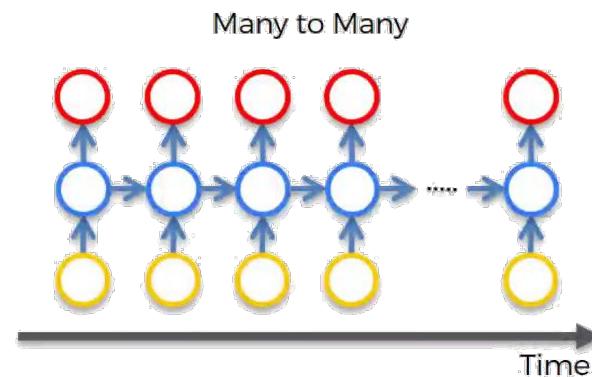
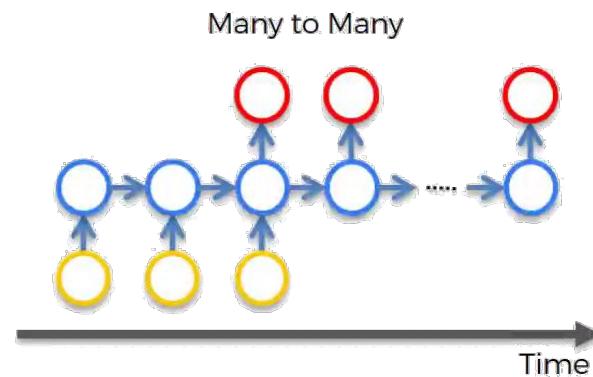


"black and white
dog jumps over
bar."



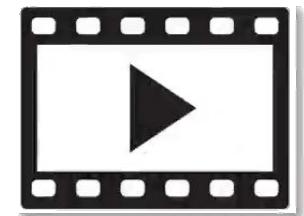
"Thanks for a great
party at the
weekend, we really
enjoyed it!"

sentiment: positive
score: 86%



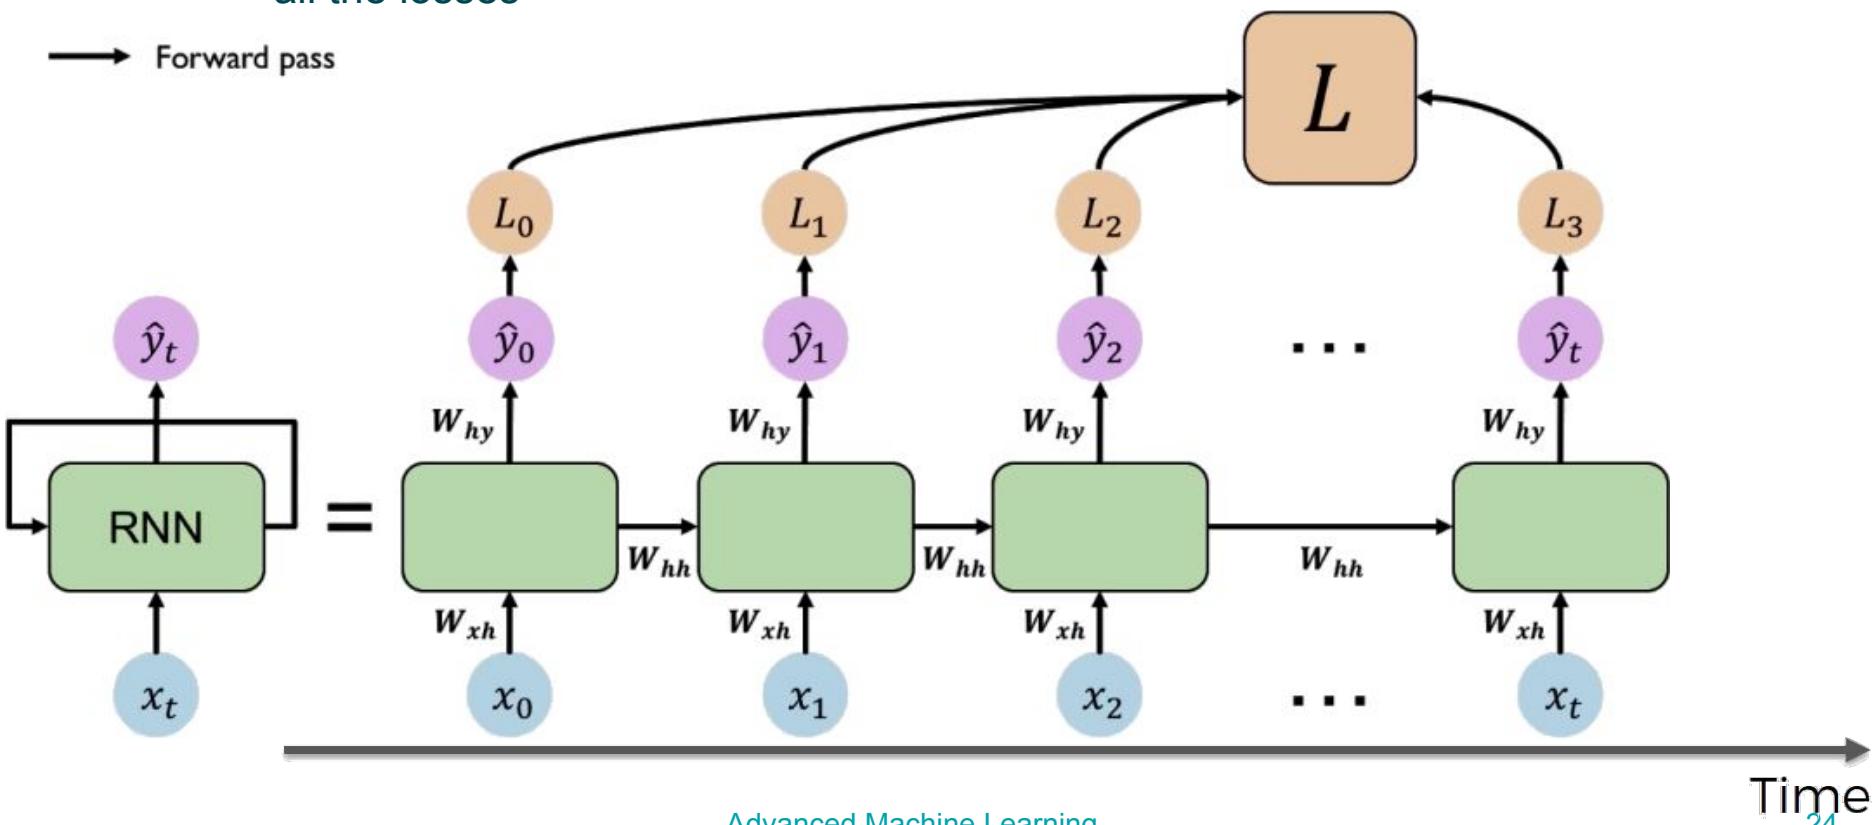
Translate into Spanish ▾

Soy un chico al que le gusta aprender



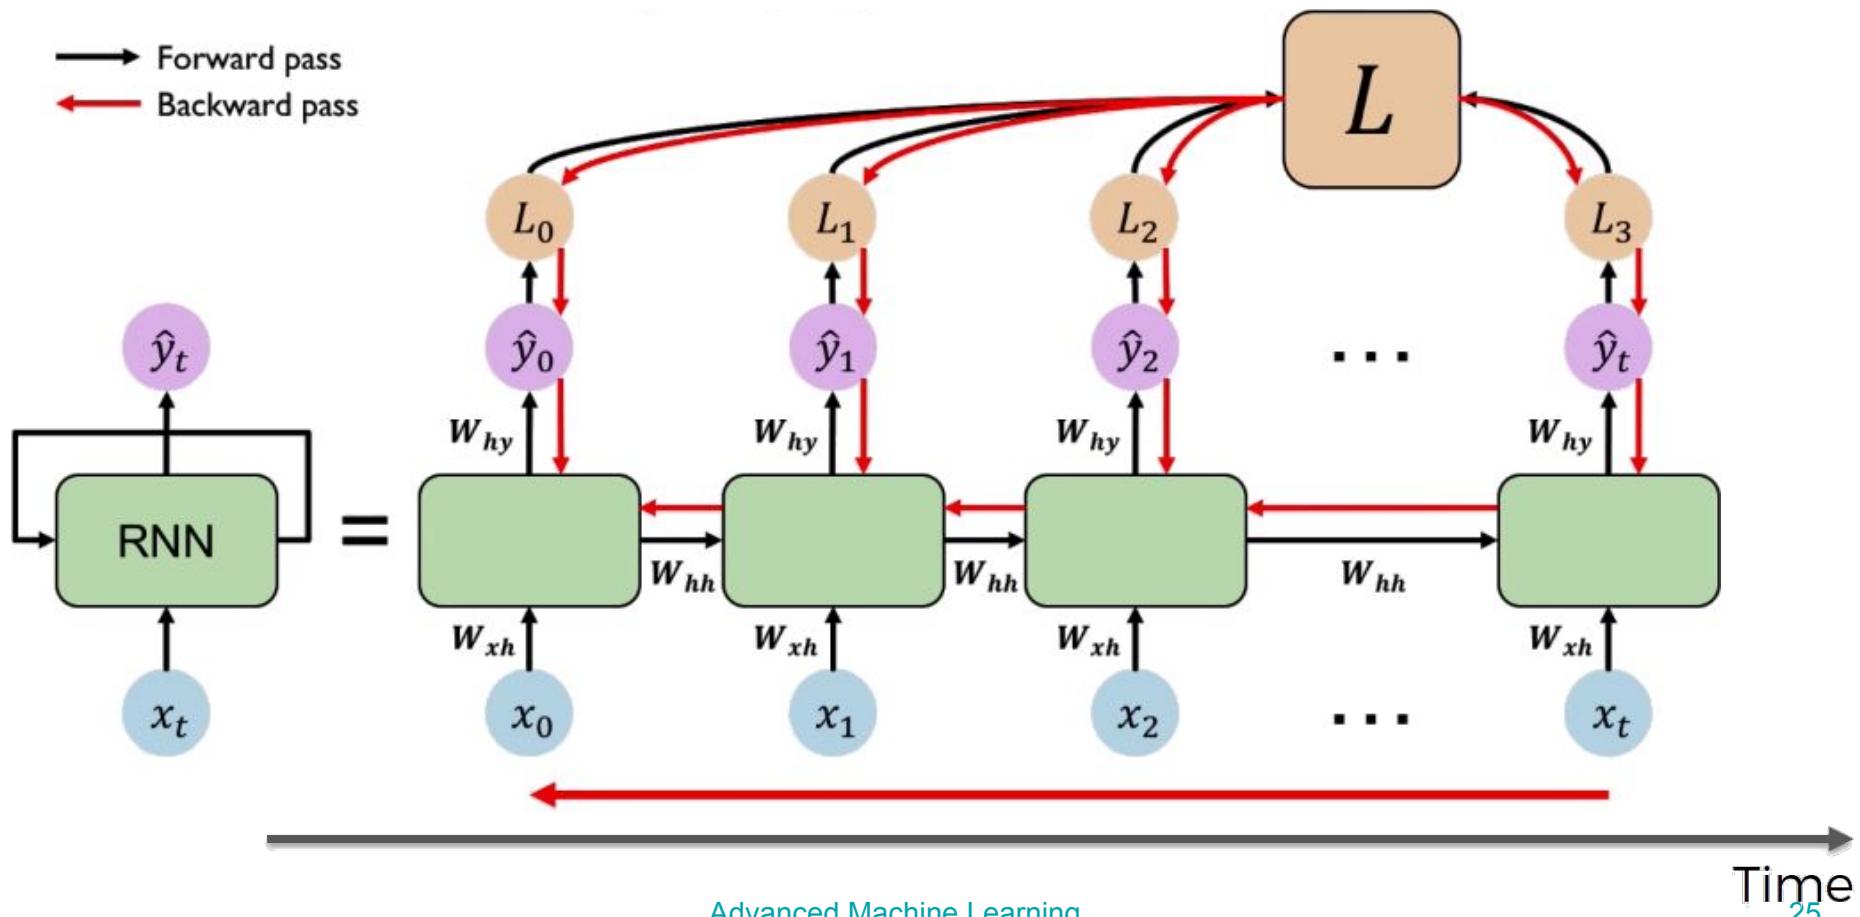
RNN, Forward pass and Loss calculation

- We can depict RNNs as:
 - Multiple copies of themselves, that pass down a message through time (h_t)
 - REMEMBER: we reuse the same Weight matrices at each time step
 - In the forward pass:
 - We can compute a loss for each time step, and the global loss will be the sum of all the losses



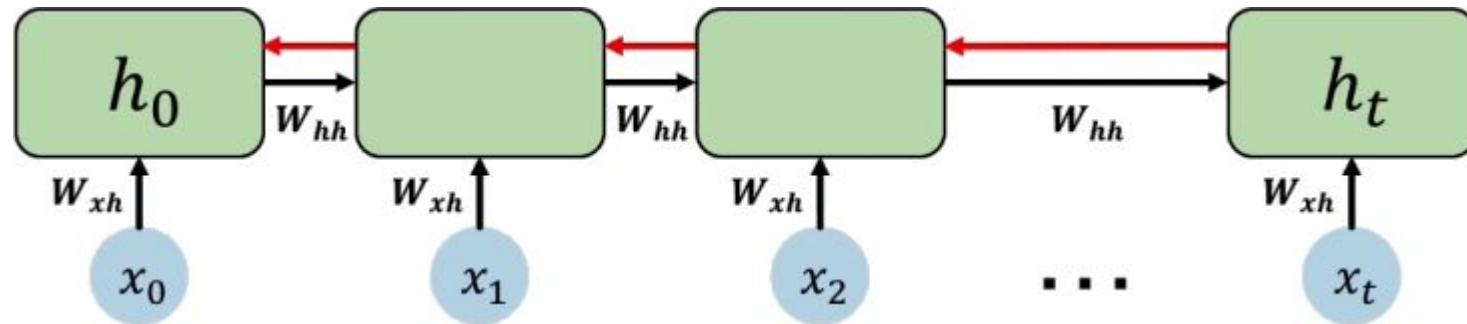
RNN-> Backpropagation Through Time

- BPTT
 - Error is Back Propagated
 - First through each state
 - Finally through Hidden states



Problems with RNNs

- The RNN : “Vanishing / Exploding Gradient”
- Computing the gradient of h_0 :
 - involves many “factors” (multiplications) of W_{hh} and activation function derivatives



Razvan Pascanu et al. (2013)

Many values > 1 :
exploding gradients

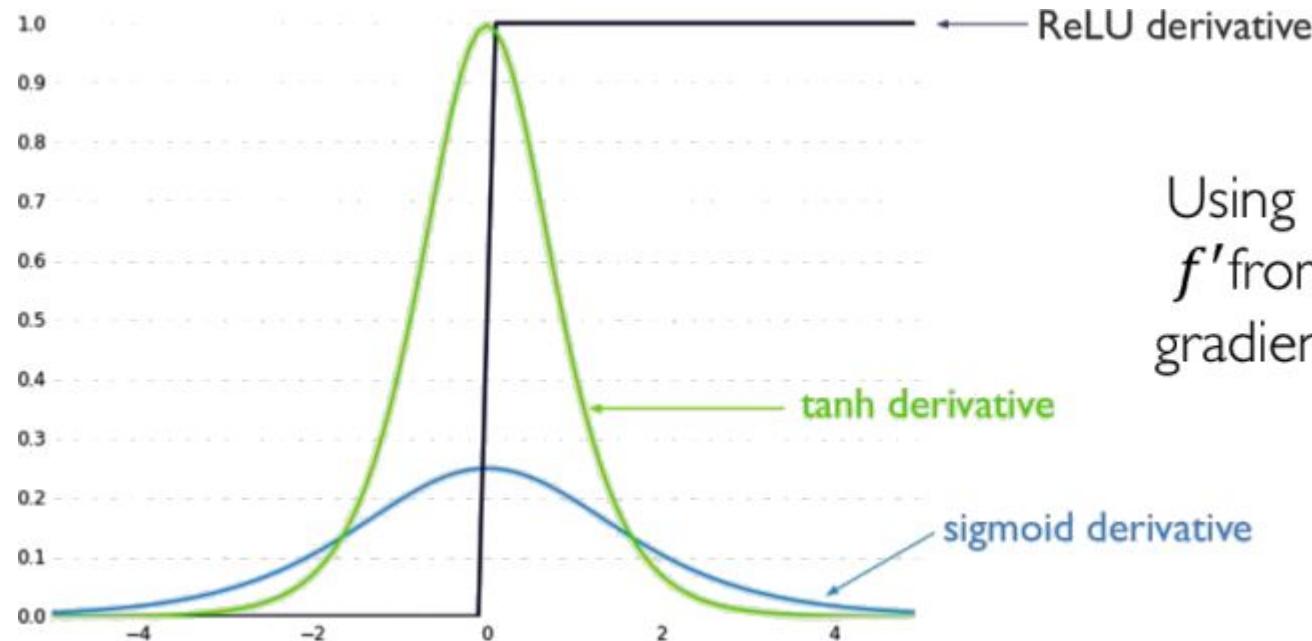
Gradient clipping to
scale big gradients

Many values < 1 :
vanishing gradients

1. Activation function
2. Weight initialization
3. Network architecture

Tackling Vanishing Gradient Problem

- Strategy 1: Choose activation function whose derivative is = 1



Using ReLU prevents f' from shrinking the gradients when $x > 0$

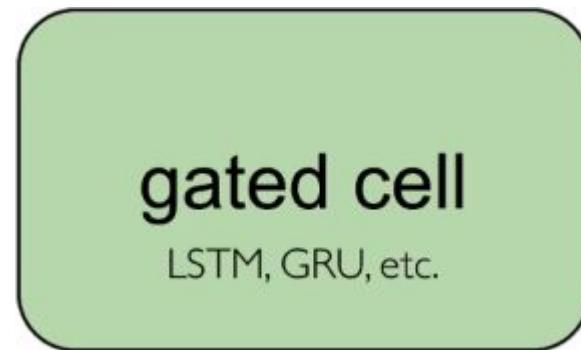
Tackling Vanishing Gradient Problem

- Strategy 2: Initialise the weights to identity matrix
 - Usually, Weights are initialized to a number close to 0
 - Initialising them to identity matrix, helps from shrinking too quickly

$$I_n = \begin{pmatrix} 1 & 0 & 0 & \cdots & 0 \\ 0 & 1 & 0 & \cdots & 0 \\ 0 & 0 & 1 & \cdots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & 0 & \cdots & 1 \end{pmatrix}$$

Tackling Vanishing Gradient Problem

- Strategy 3: **Using complex recurrent units with gates**
 - Instead of using vanilla RNN
 - These help control the information that is passing to the future

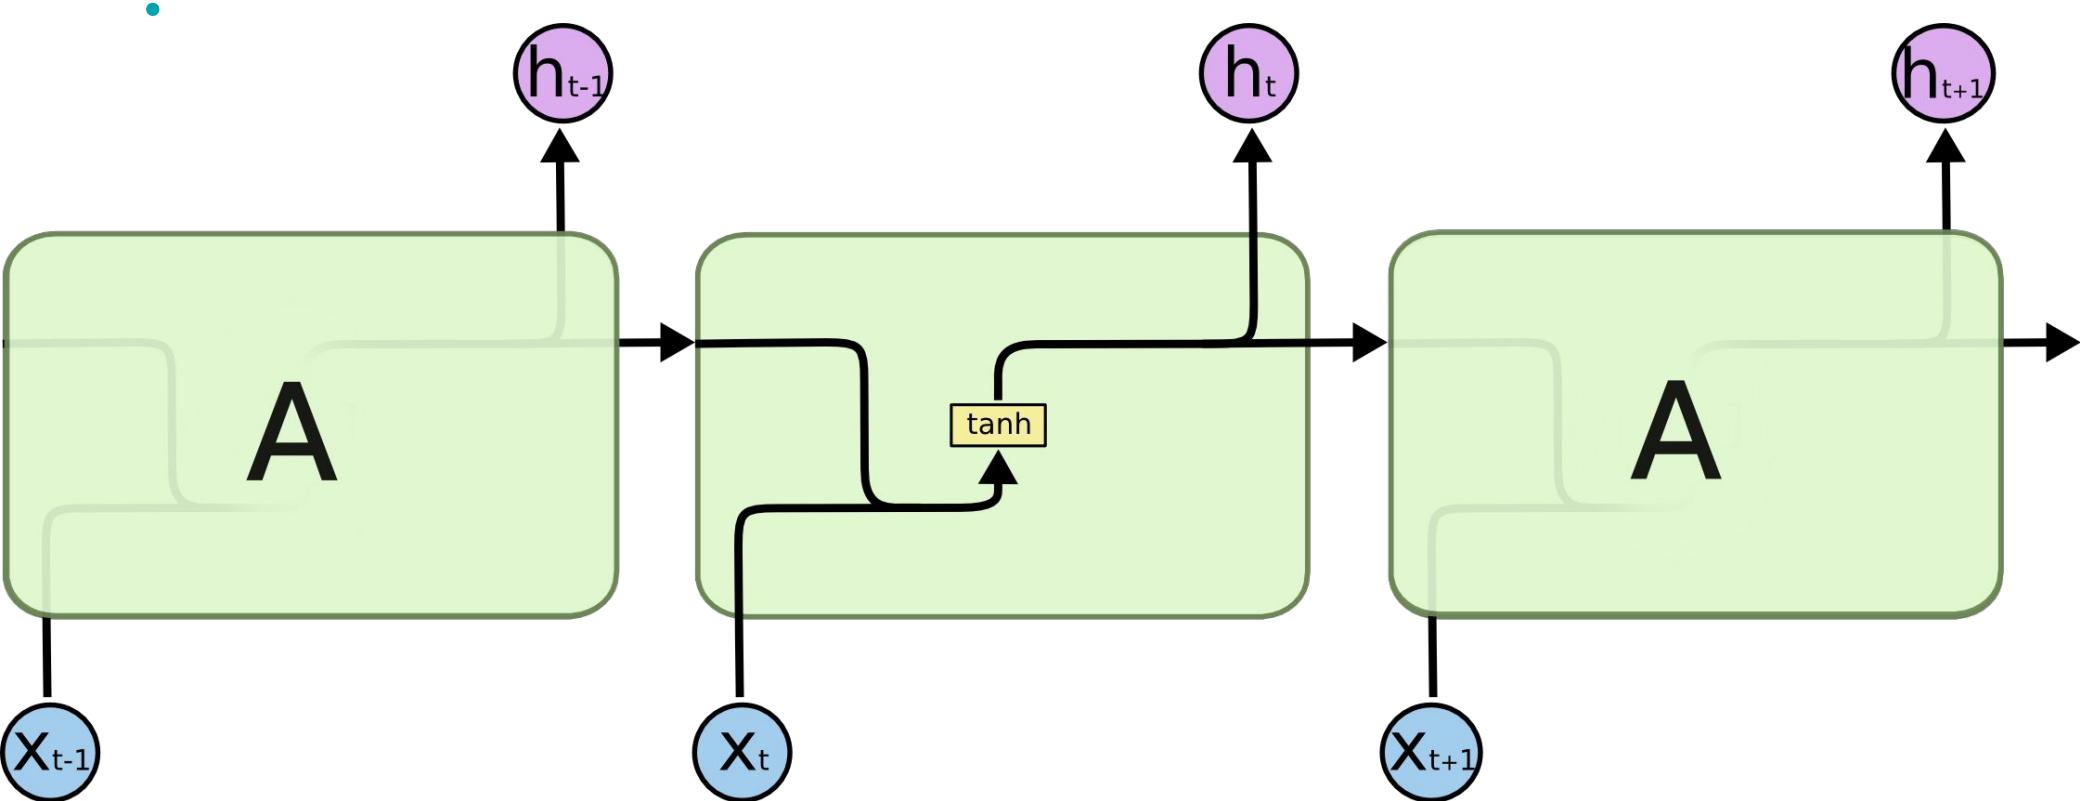


Long Short-Time Memory Networks (LSTM)

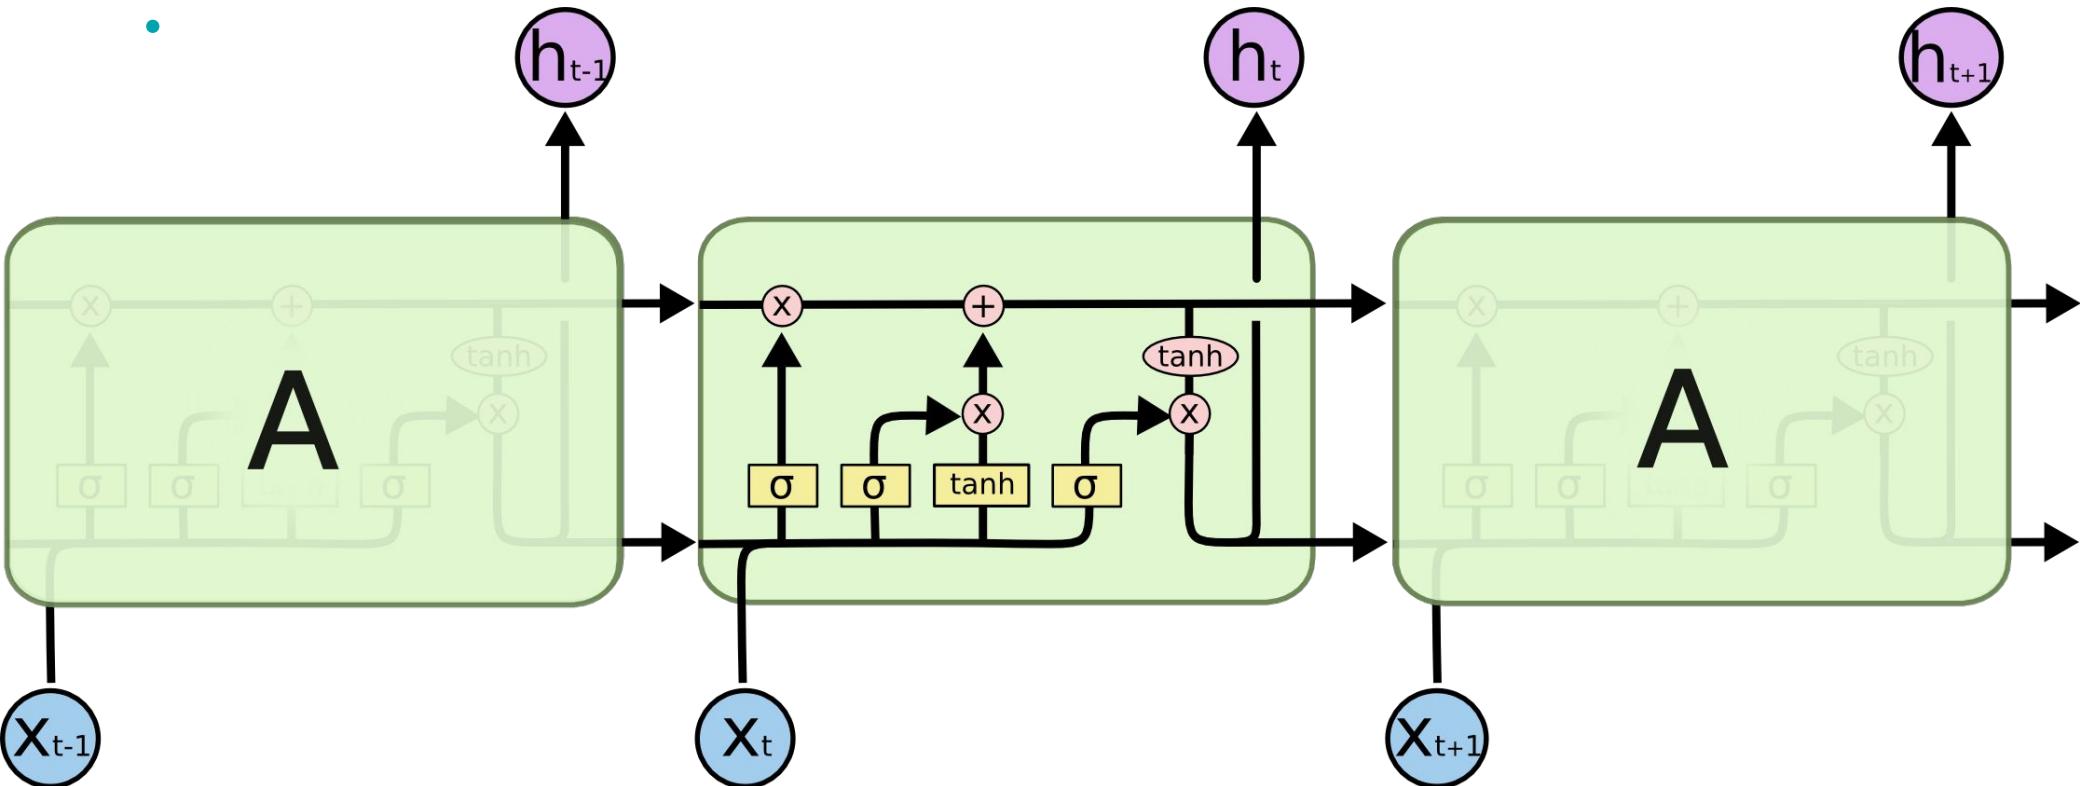
```
tf.keras.layers.LSTM(units)
```



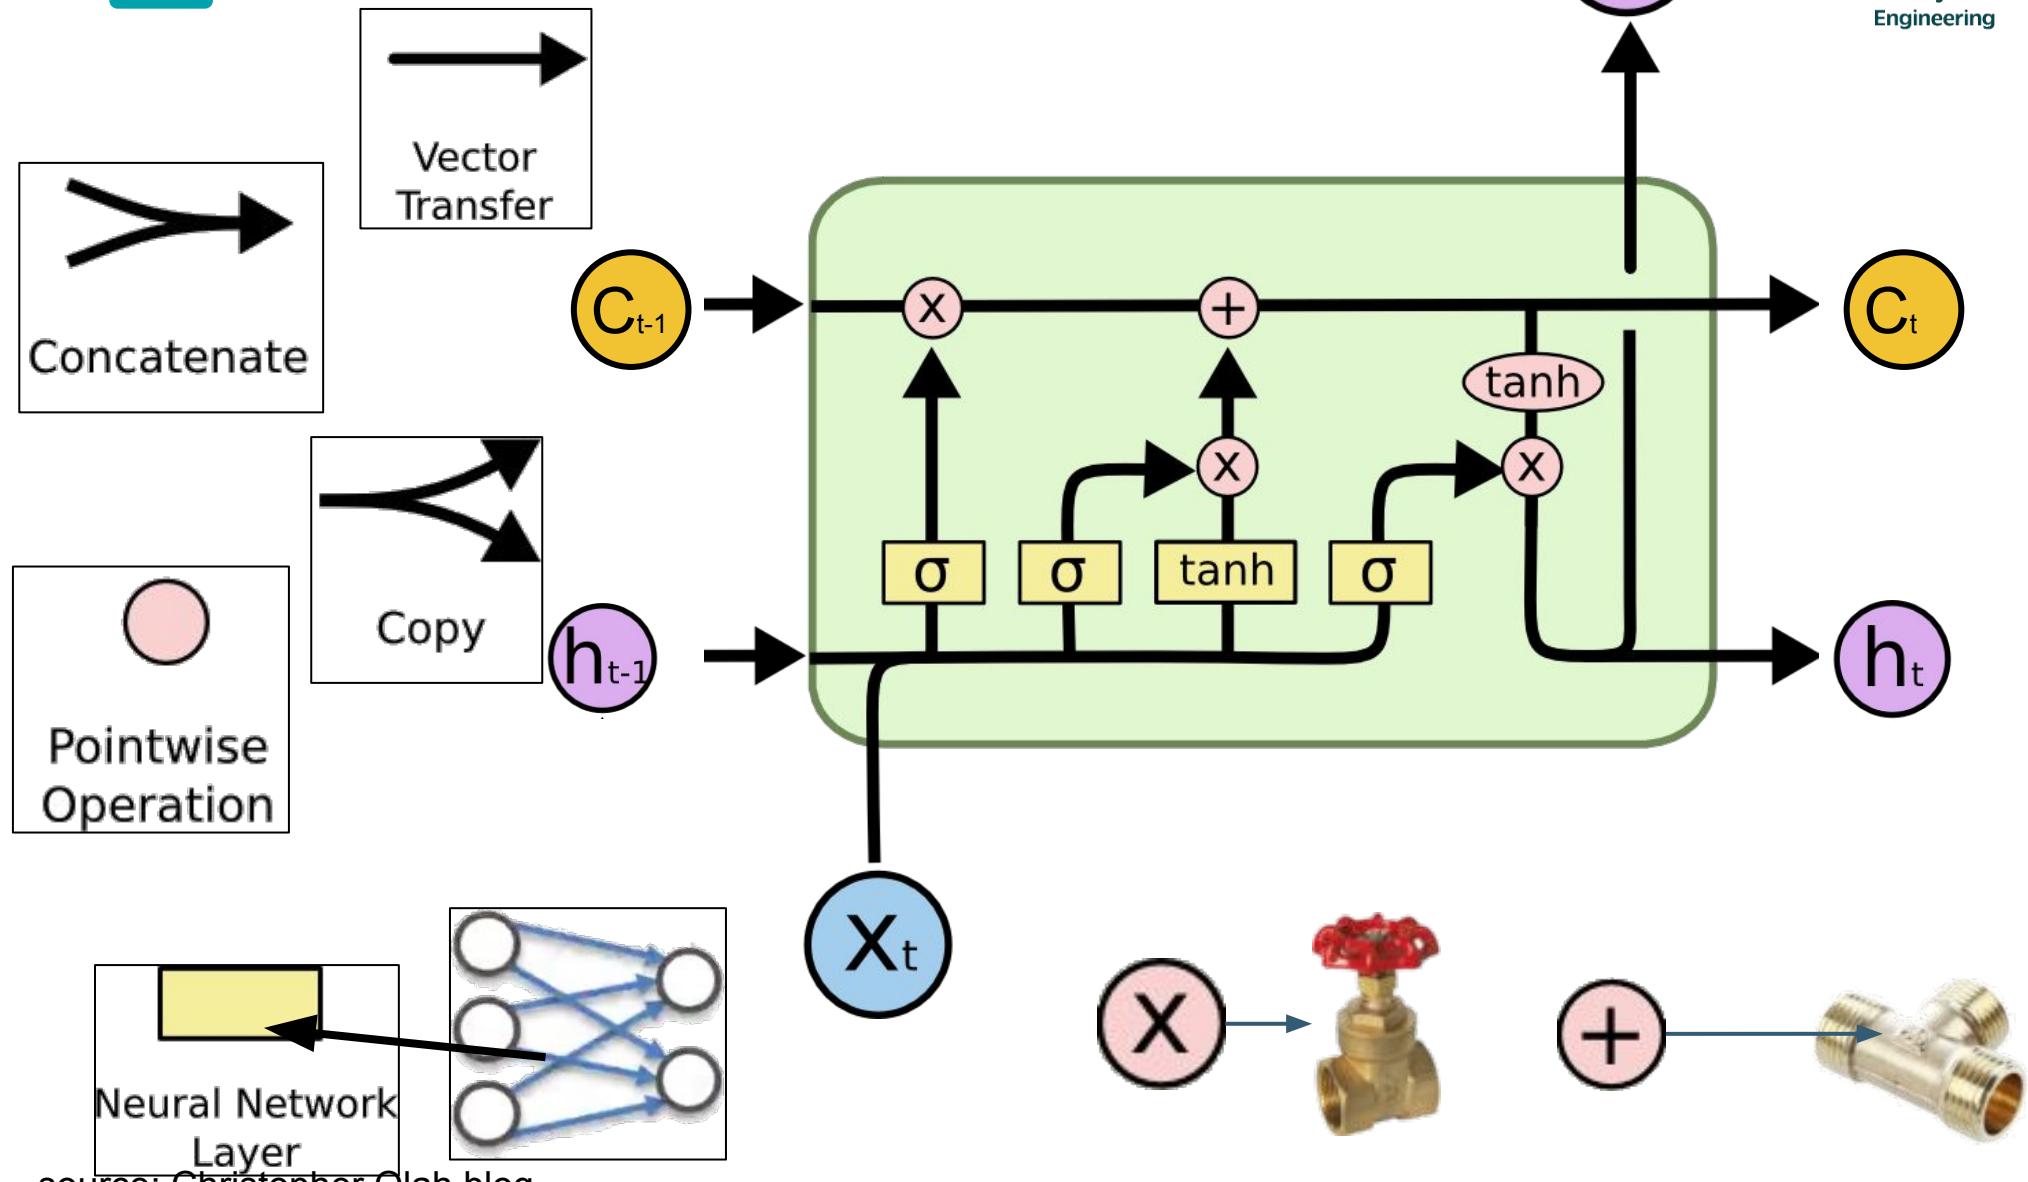
Normal RNN



LSTM - simplified

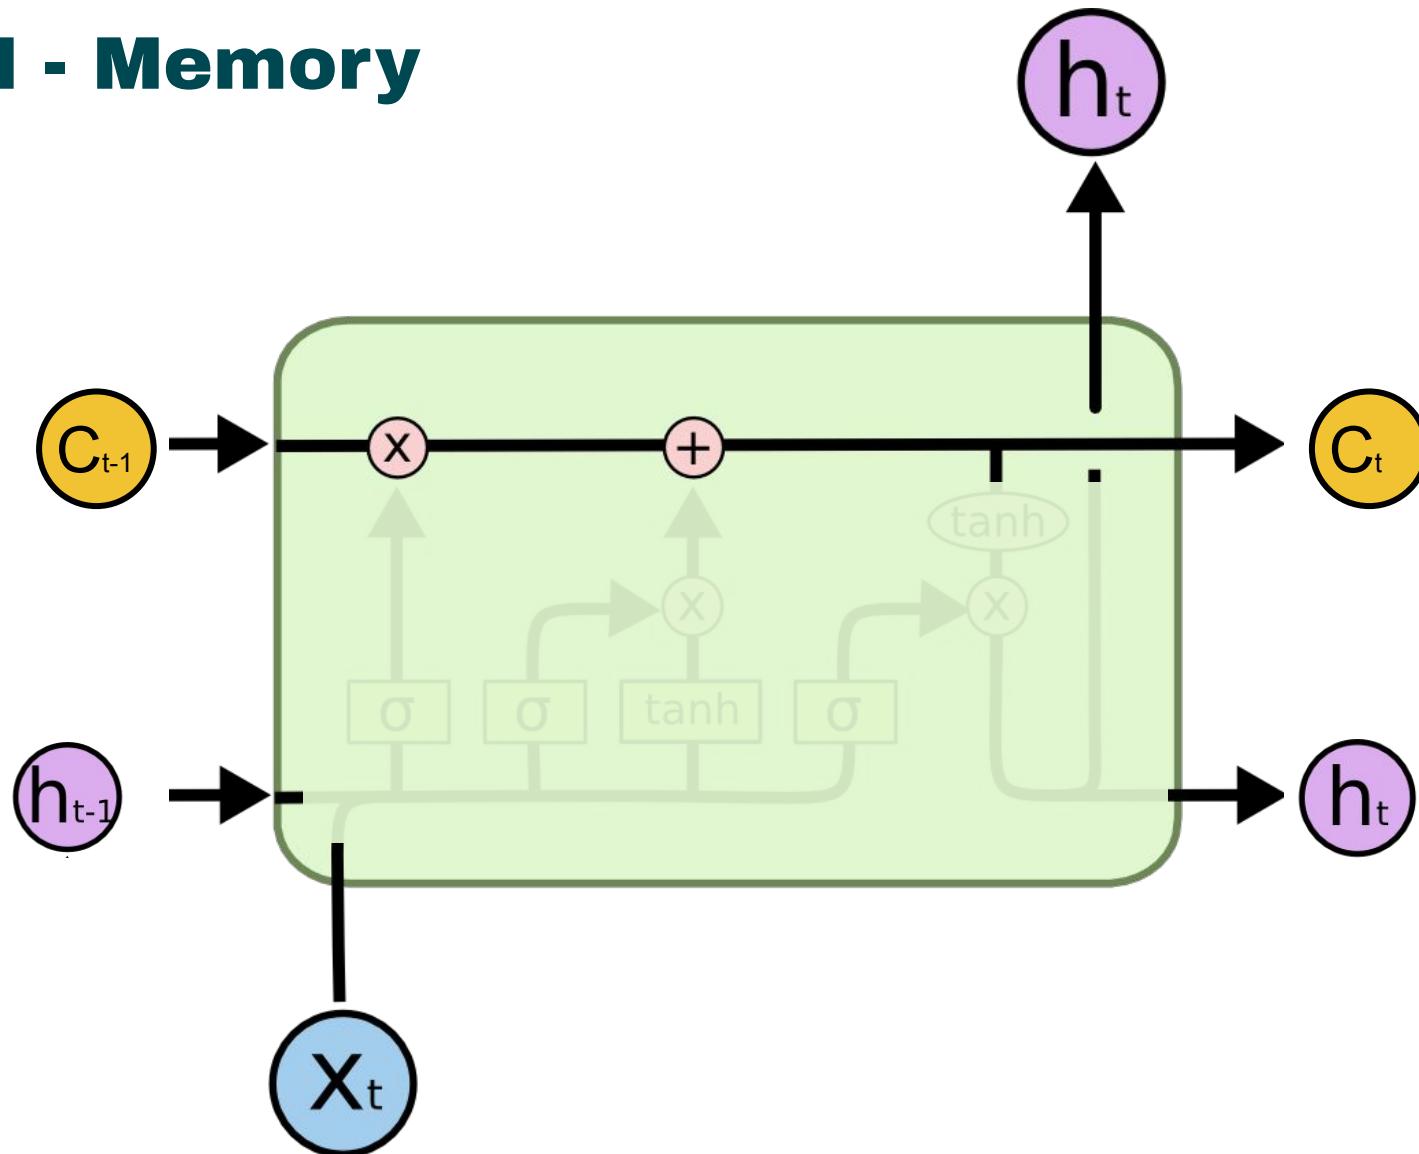


LSTM - simplified

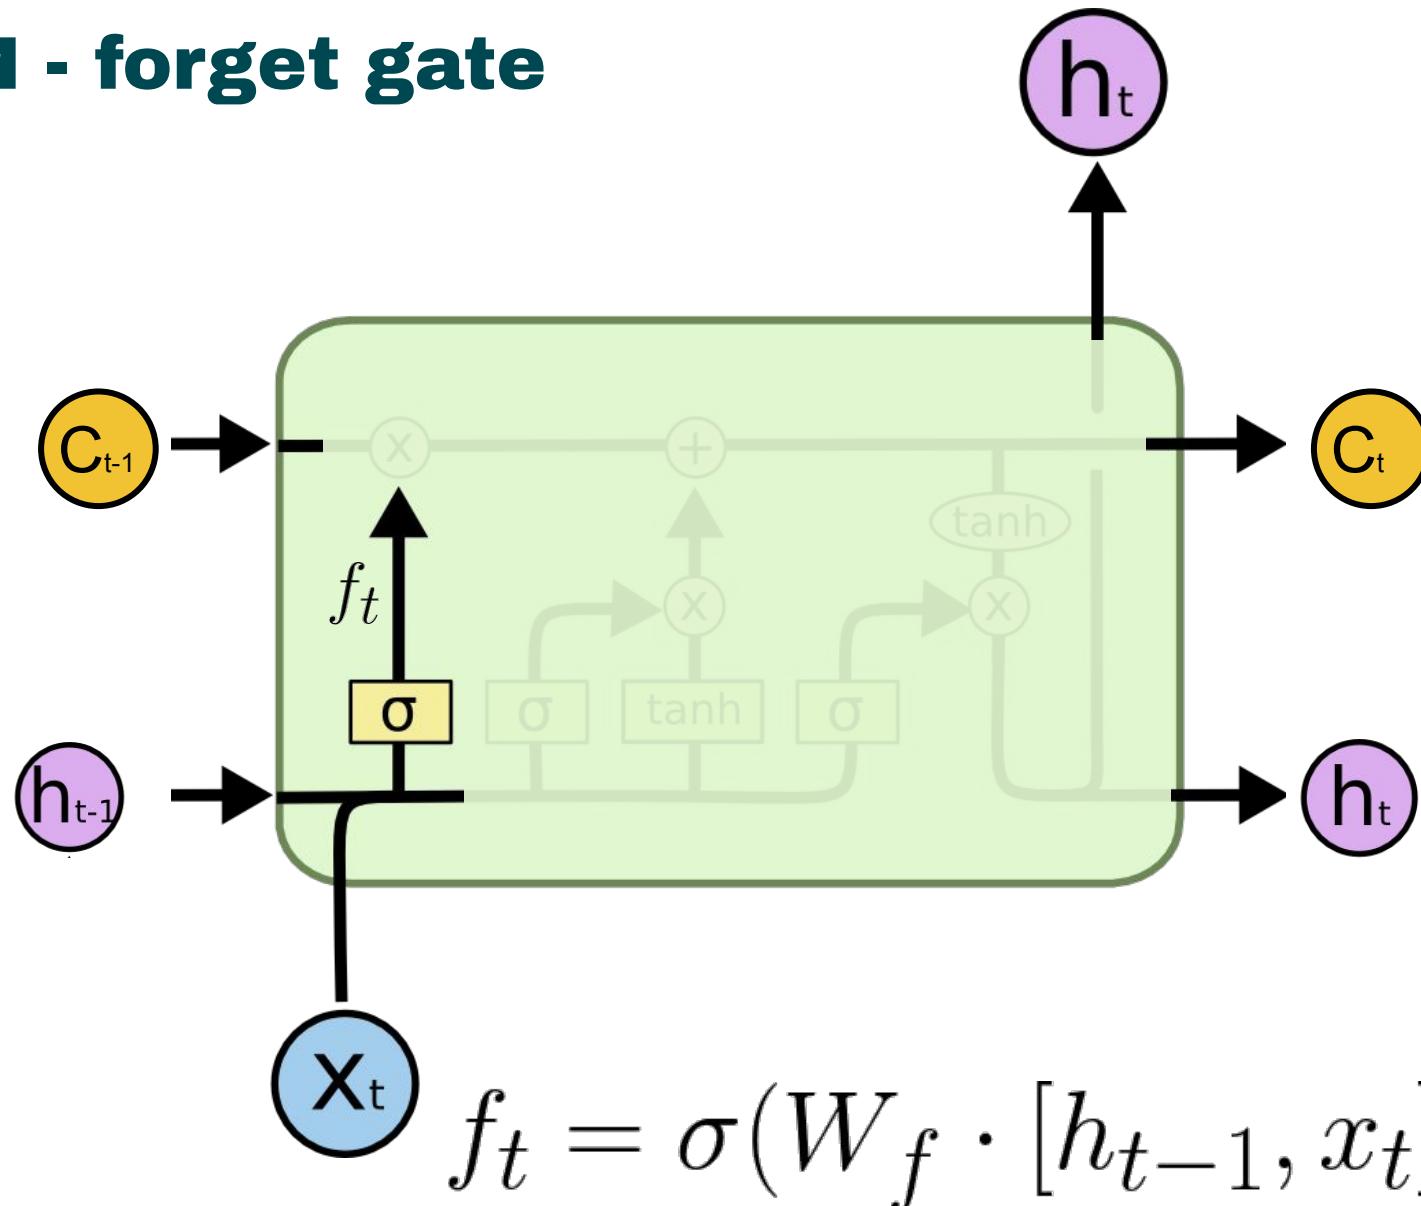


source: Christopher Olah blog

LSTM - Memory

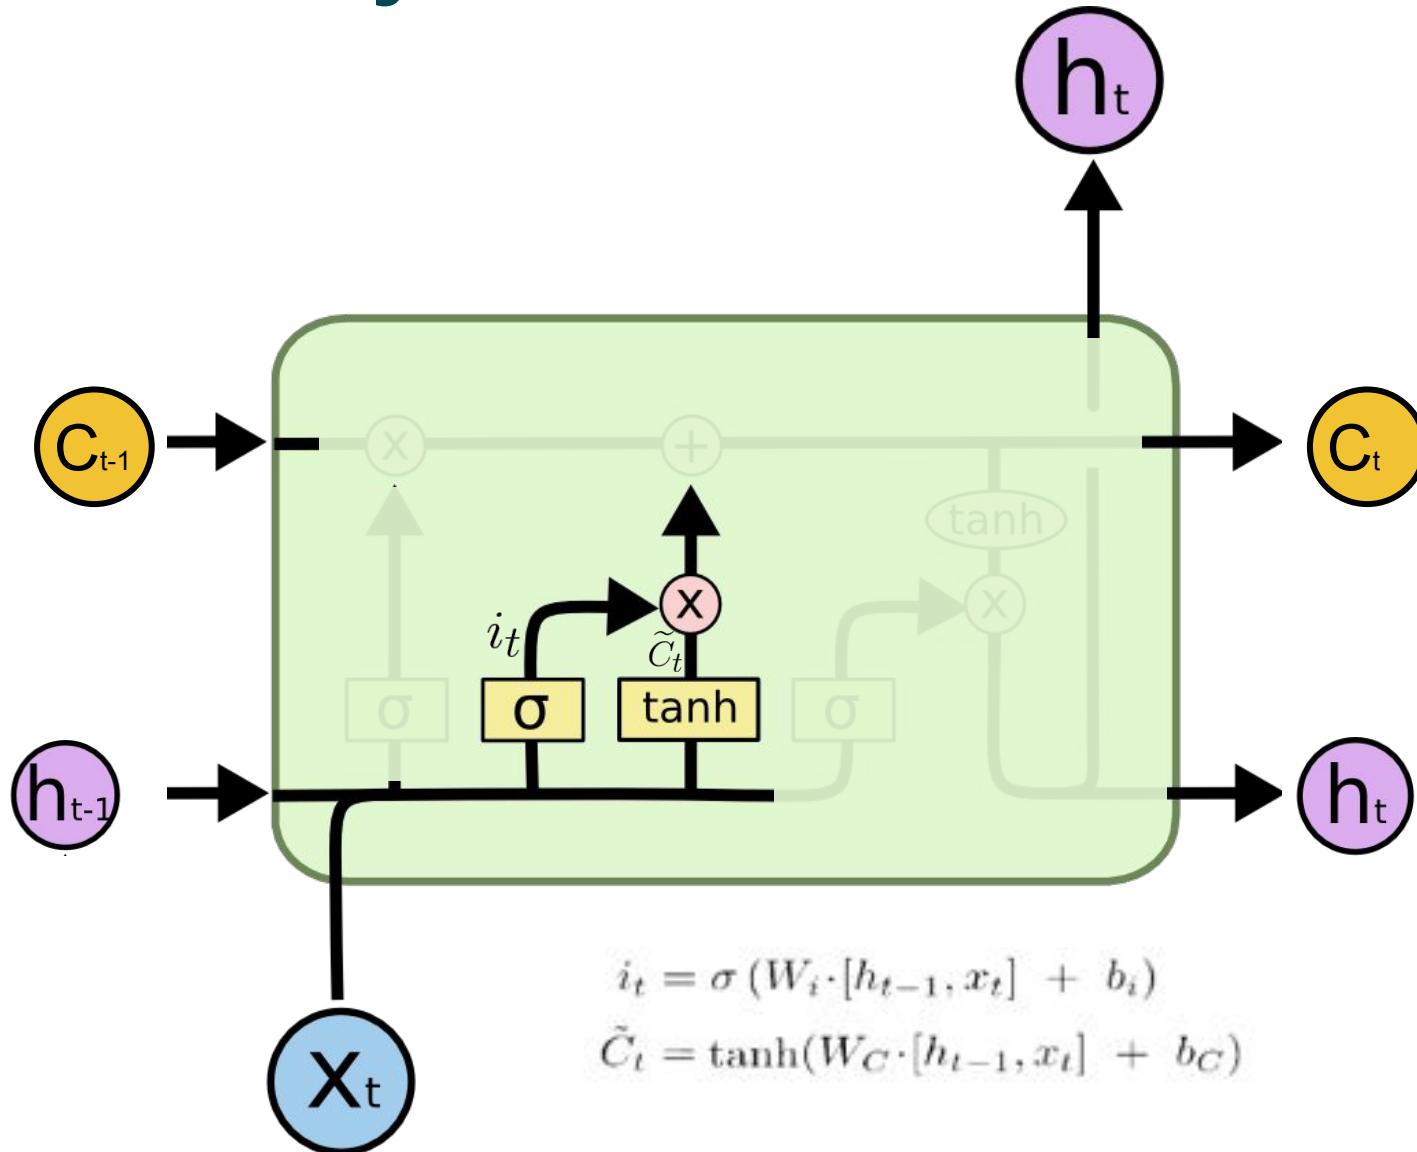


LSTM - forget gate

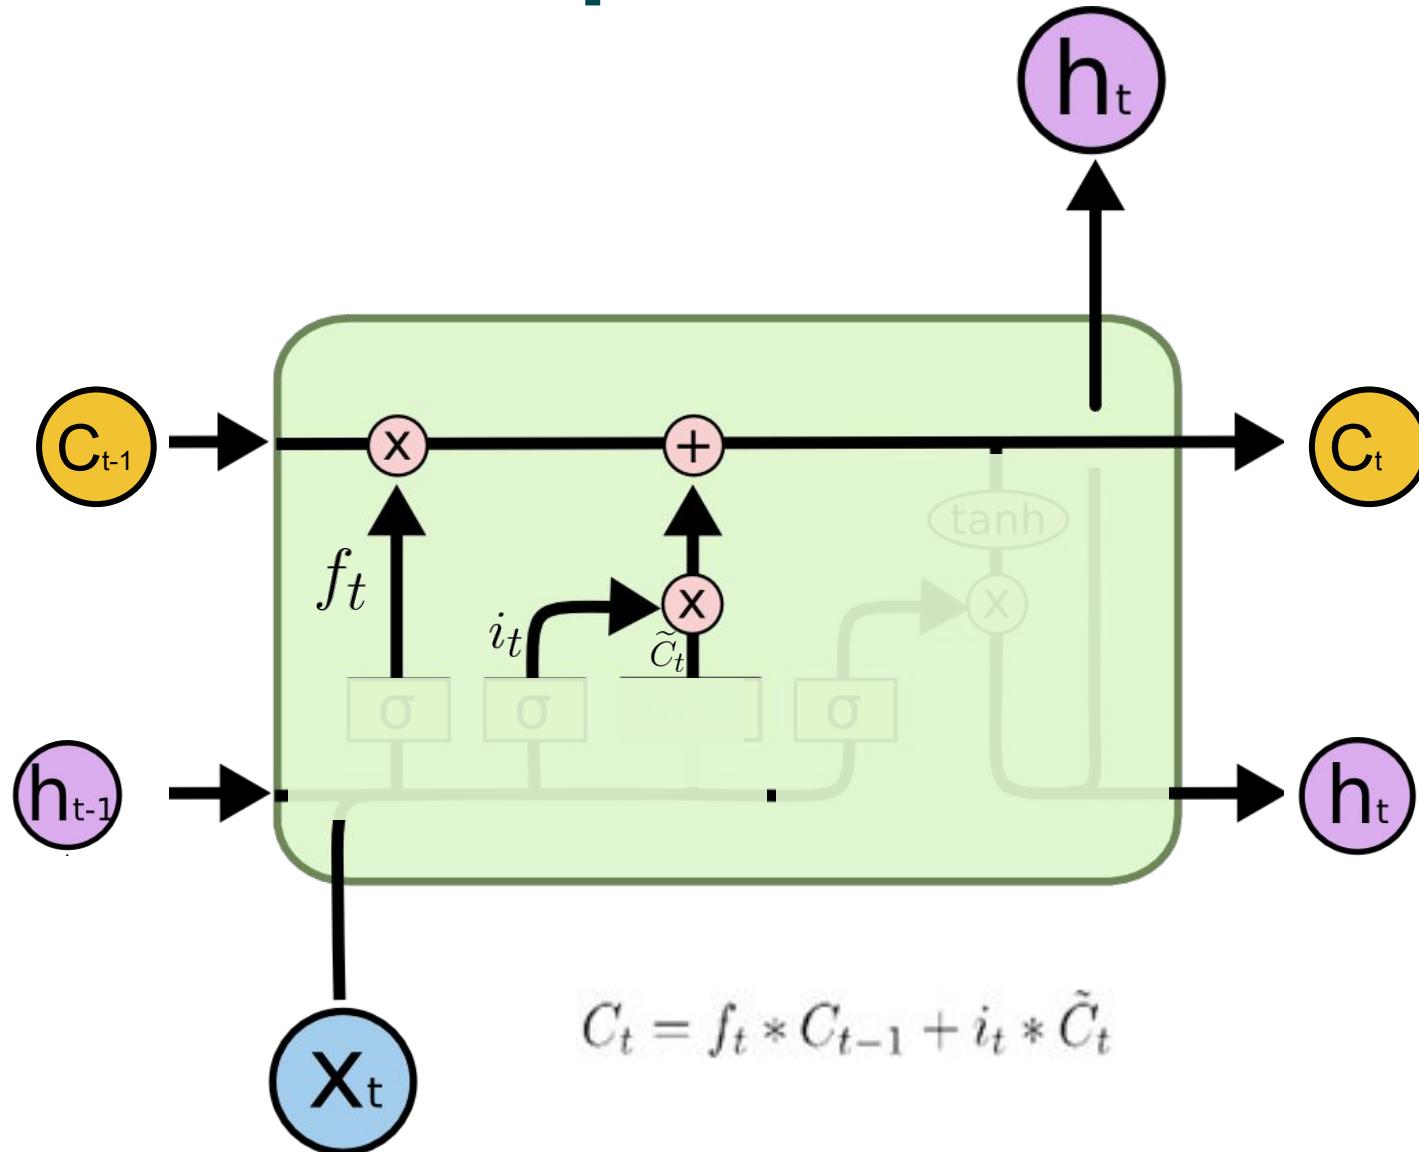


$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$

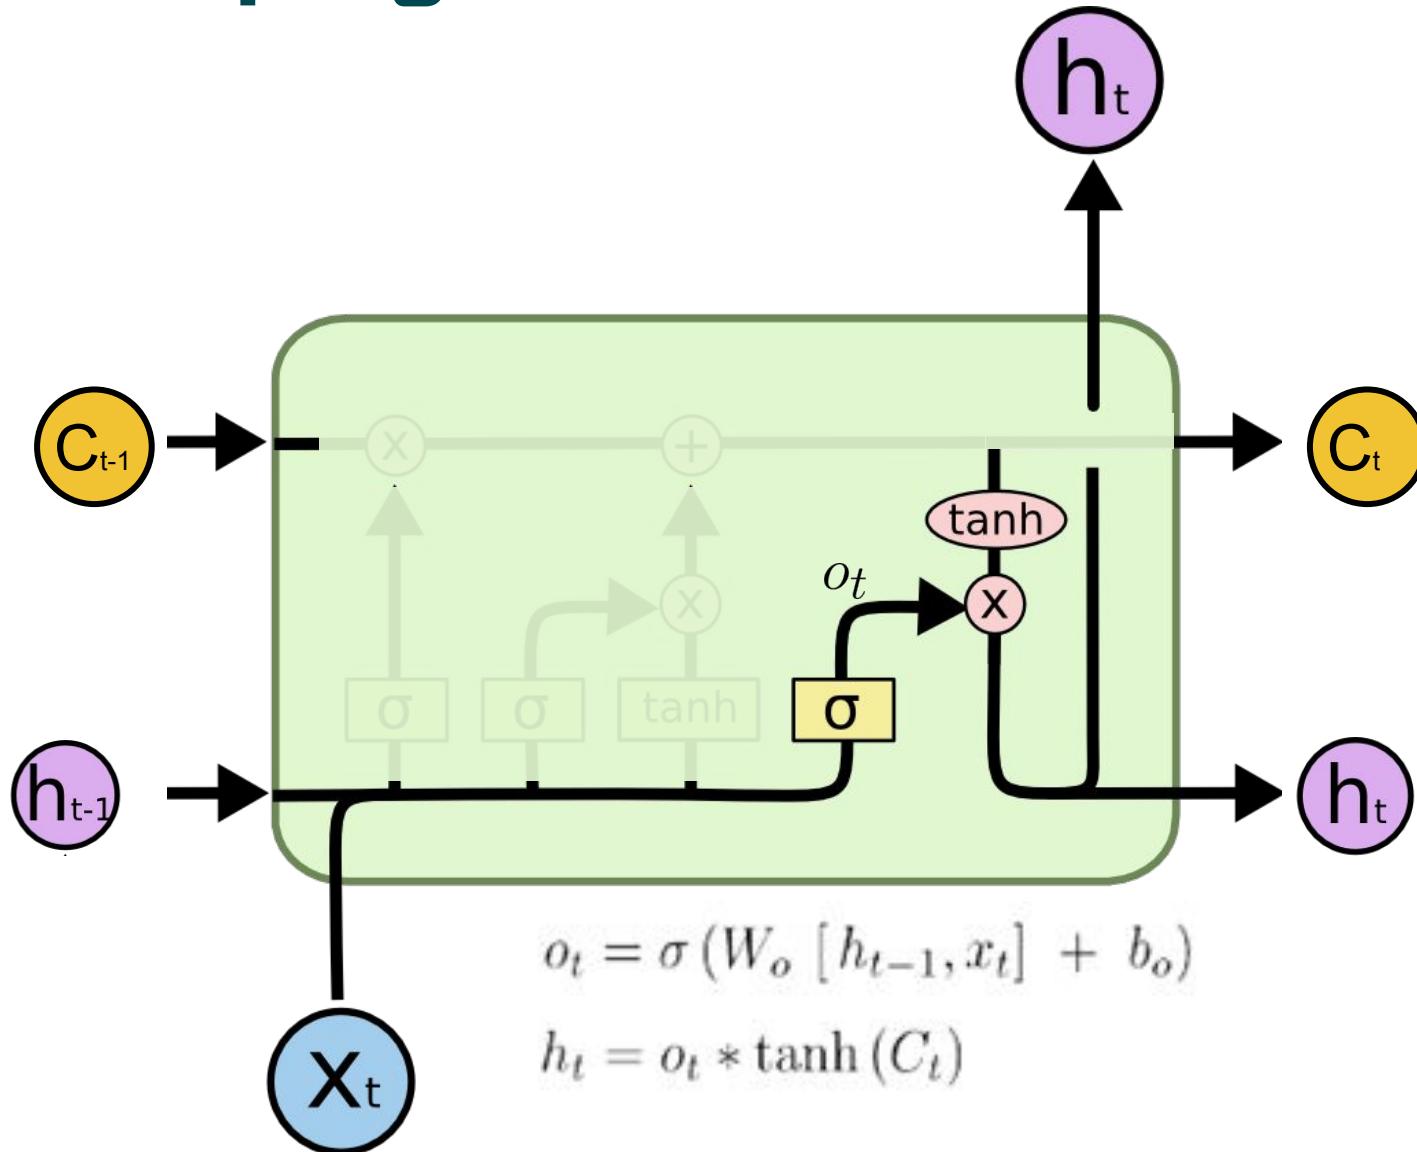
LSTM - Memory Gate



LSTM - Cell State Update



LSTM - Output gate



LSTM - Cosas a tener en cuenta

- El número de variables a optimizar es muy grande por cada LSTM:
 - pesos y bias del forget gate
 - pesos y bias del memory gate X2
 - Pesos y bias del output gate
- Todo eso, por cada neurona LSTM!

Son muy potentes, pero, el tiempo de aprendizaje es alto, hay que diseñarlos con cabeza, si no, perdemos mucho tiempo.

Gated Recurrent Unit (GRU)

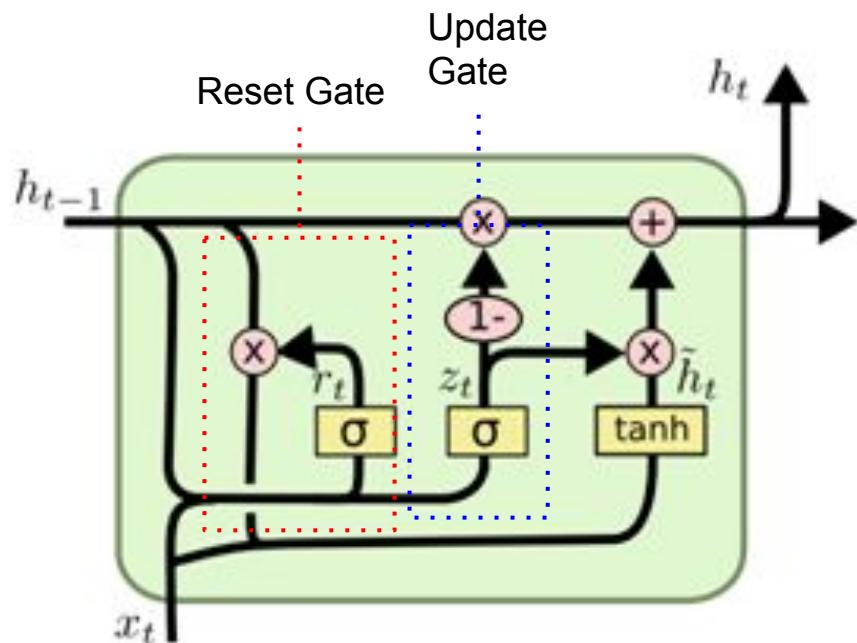
```
tf.keras.layers.GRU(units)
```



RNN - GRU

- Similar to LSTM
 - They combine the forget and input gate into only one gate, called the update gate
 - They also merge the Cell state and the hidden state into one

RNN - Gated Recurrent Unit



$$z_t = \sigma (W_z \cdot [h_{t-1}, x_t])$$

$$r_t = \sigma (W_r \cdot [h_{t-1}, x_t])$$

$$\tilde{h}_t = \tanh (W \cdot [r_t * h_{t-1}, x_t])$$

$$h_t = (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t$$

RNN - GRU

- The main advantage of GRU is **the reduced number of parameters** as compared to LSTM's without any compromise whatsoever which has resulted in **faster convergence** and a **more generalized model**.

Advanced use of RNNs

RNN - Recurrent Dropout

- RNNs also tend to Overfit.
- But applying dropout is not trivial for RNN
- The same pattern of dropped units should be applied every time step
 - Otherwise instead of helping learn you achieve the contrary (*Yarin Gal 2015*)
- Keras has 2 dropouts for RNNs
 - Dropout-> specifies dropout rate for input units of the layer
 - Recurrent_dropout-> specifies dropout rate for RNN units



```
model.add(layers.GRU(64, activation='relu',  
                     dropout=0.1,  
                     recurrent_dropout=0.5))
```

RNN - Stacking Recurrent Layers

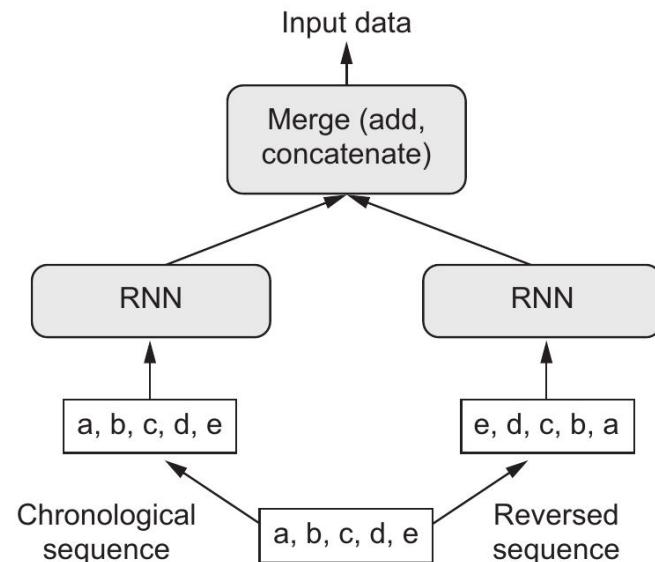
- We can stack several recurrent layers one after another
 - For example, in 2018 7 huge LSTM layers powered google translate
- In order to be able to do this, we need all intermediate layers to return their full sequence of outputs (a 3D tensor)
- This is done the next way:



```
model.add(layers.GRU(32,  
                      dropout=0.1,  
                      recurrent_dropout=0.5,  
                      return_sequences=True,  
                      input_shape=(None, float_data.shape[-1])))
```

RNN - Bidirectional RNNs

- The basic idea behind the bidirectional RNNs is to have 2 RNNs
 - One looking at the data in the given order (chronologically)
 - The other having the data reversed (anti-chronologically)
- Later, the both representations are merged
 - Helps to catch patterns that may have been overlooked

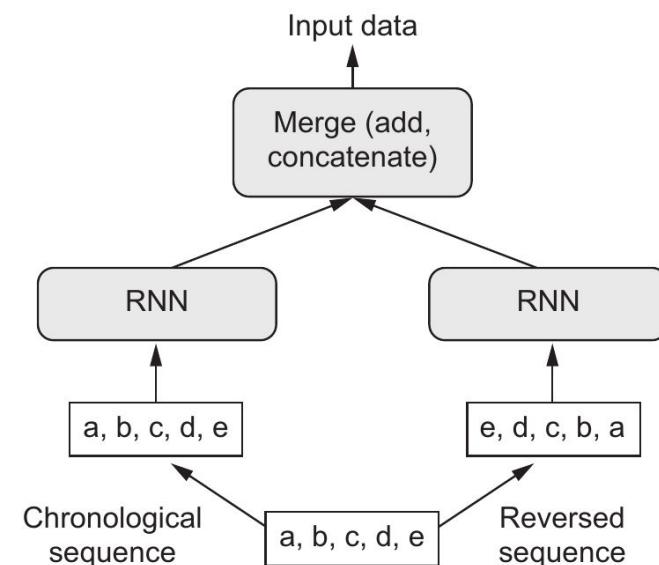


```
model = Sequential()
model.add(layers.Bidirectional(
    layers.GRU(32), input_shape=(None, float_data.shape[-1])))
```



RNN - Bidirectional RNNs

- The basic idea behind the bidirectional RNNs is to have 2 RNNs
 - One looking at the data in the given order (chronologically)
 - The other having the data reversed (anti-chronologically)
- Later, the both representations are merged
 - Helps to catch patterns that may have been overlooked



```
model = Sequential()
model.add(layers.Bidirectional(
    layers.GRU(32), input_shape=(None, float_data.shape[-1])))
```

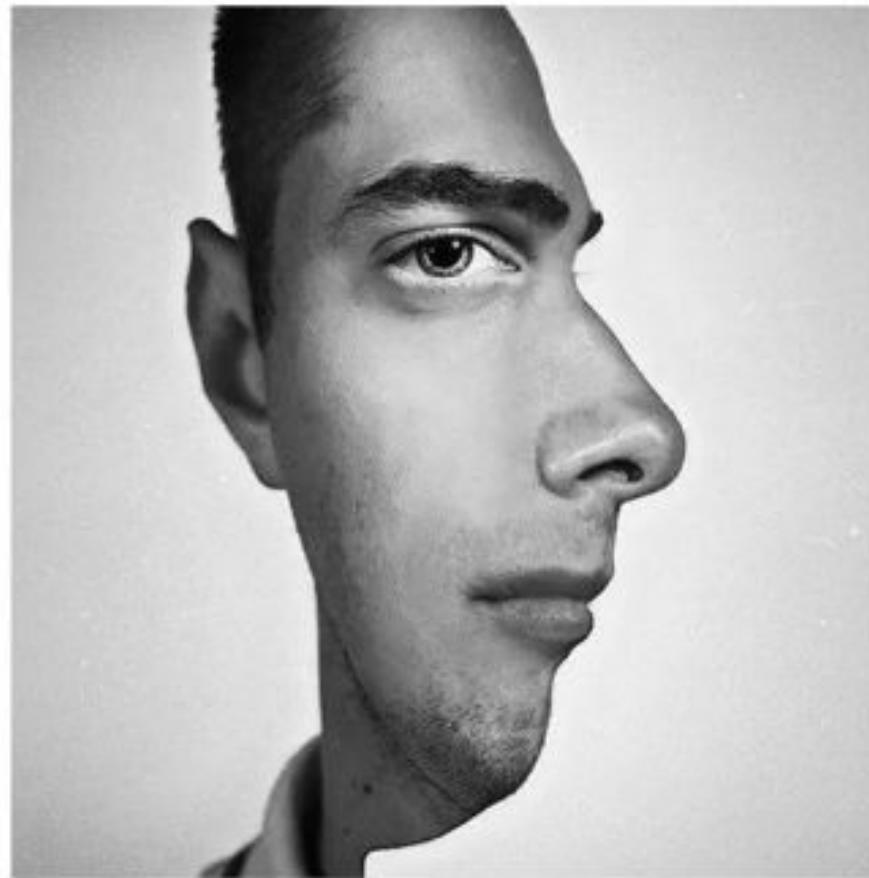


Exercises

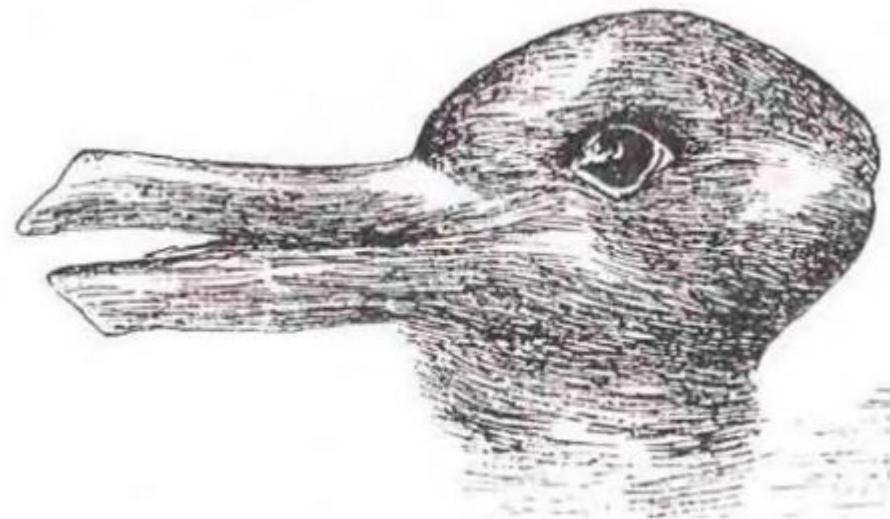
- 2 rnn examples
- exercise 1
 - google stock prediction
- exercise 2
 - text generation with rnn, moby dick

Convolutional Neural Networks

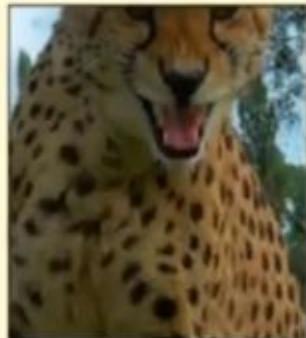
Features from images



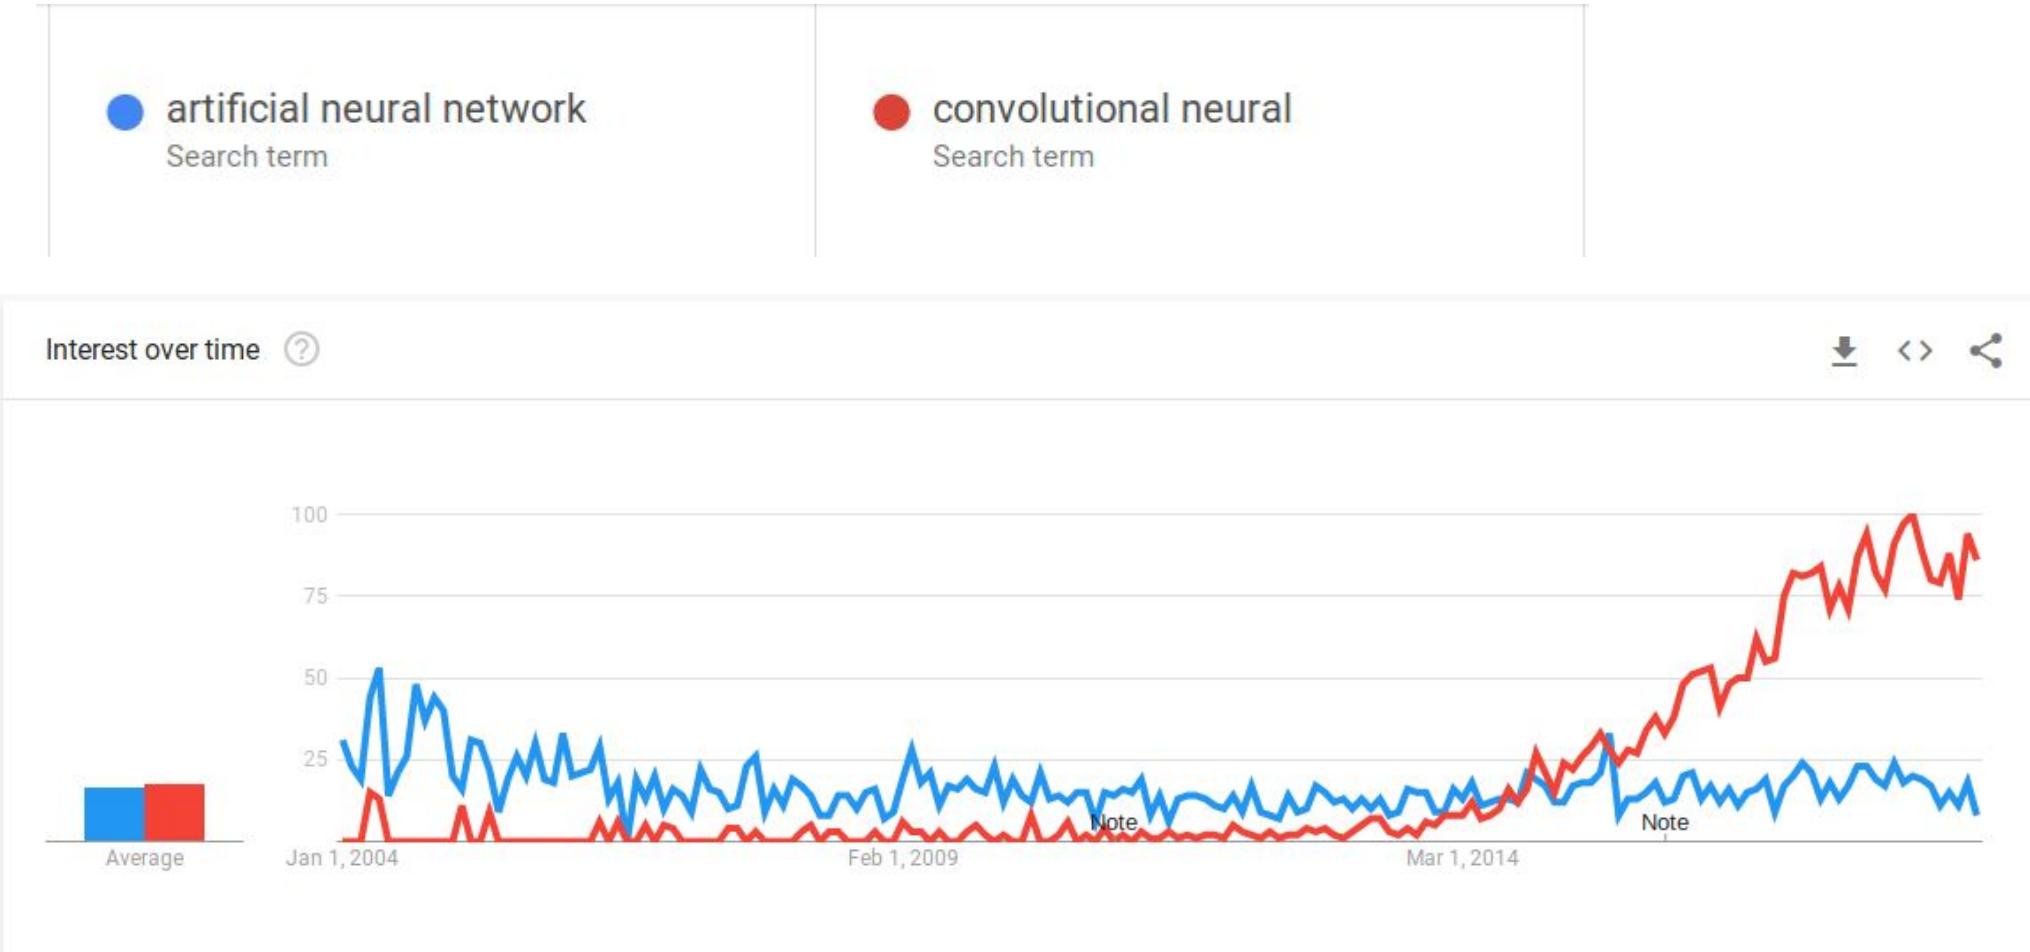
Features from images



Examples from the test set (with the network's guesses)



CNNs



CNNs

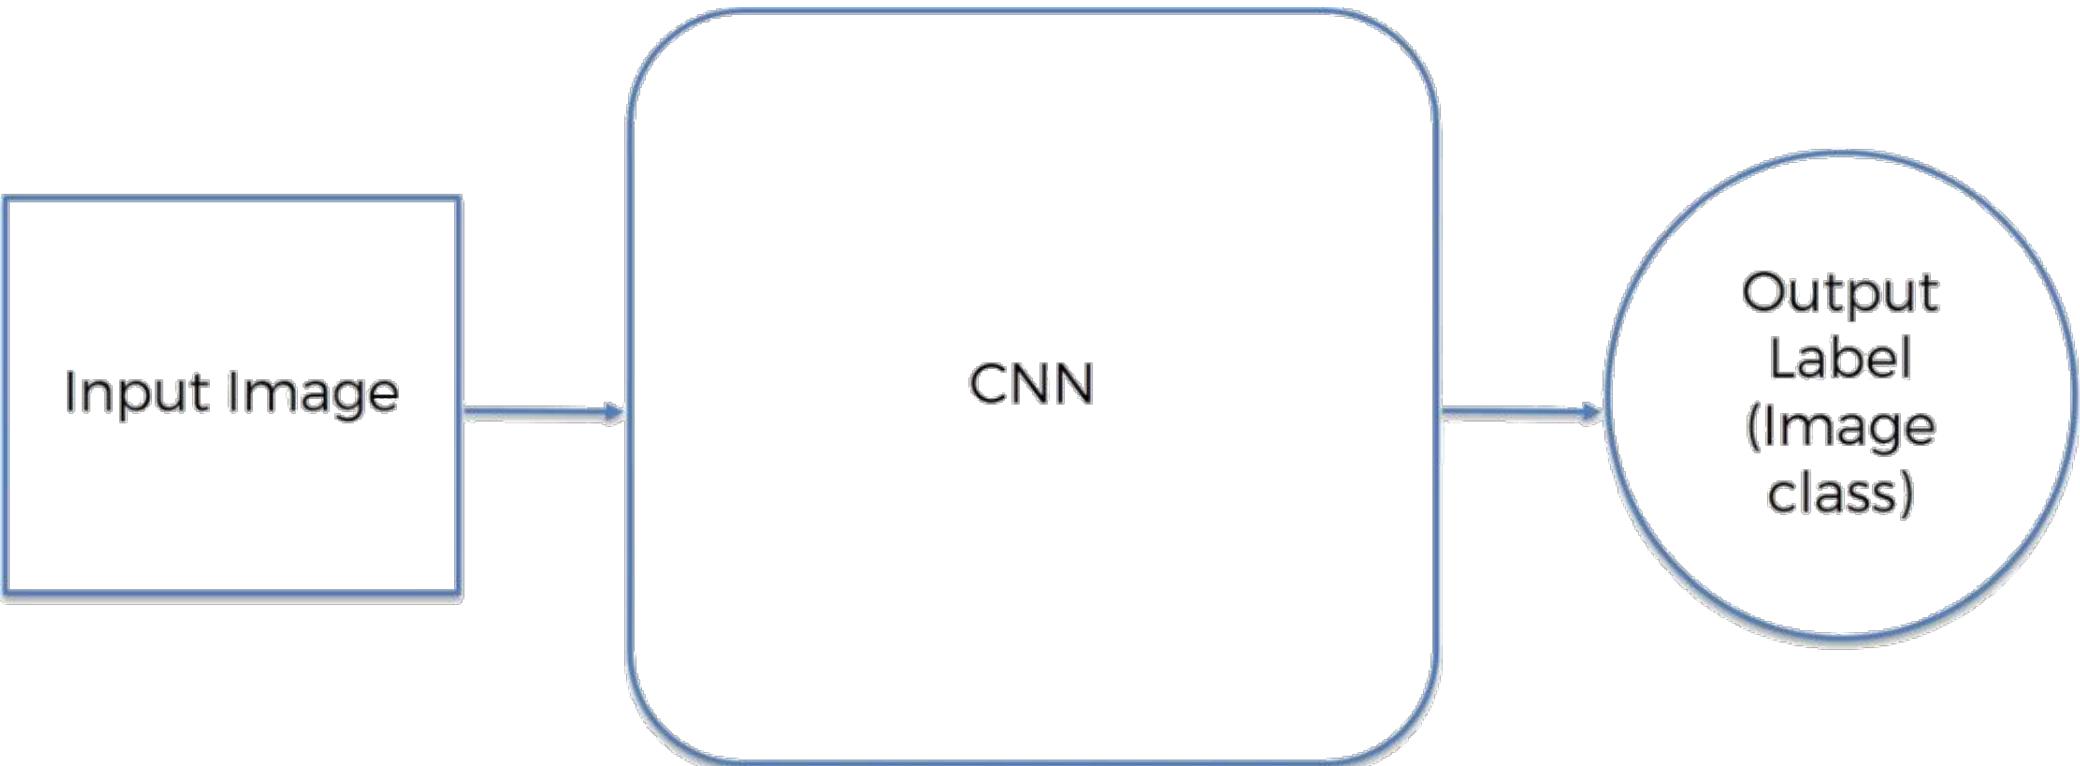


Yann Lecun

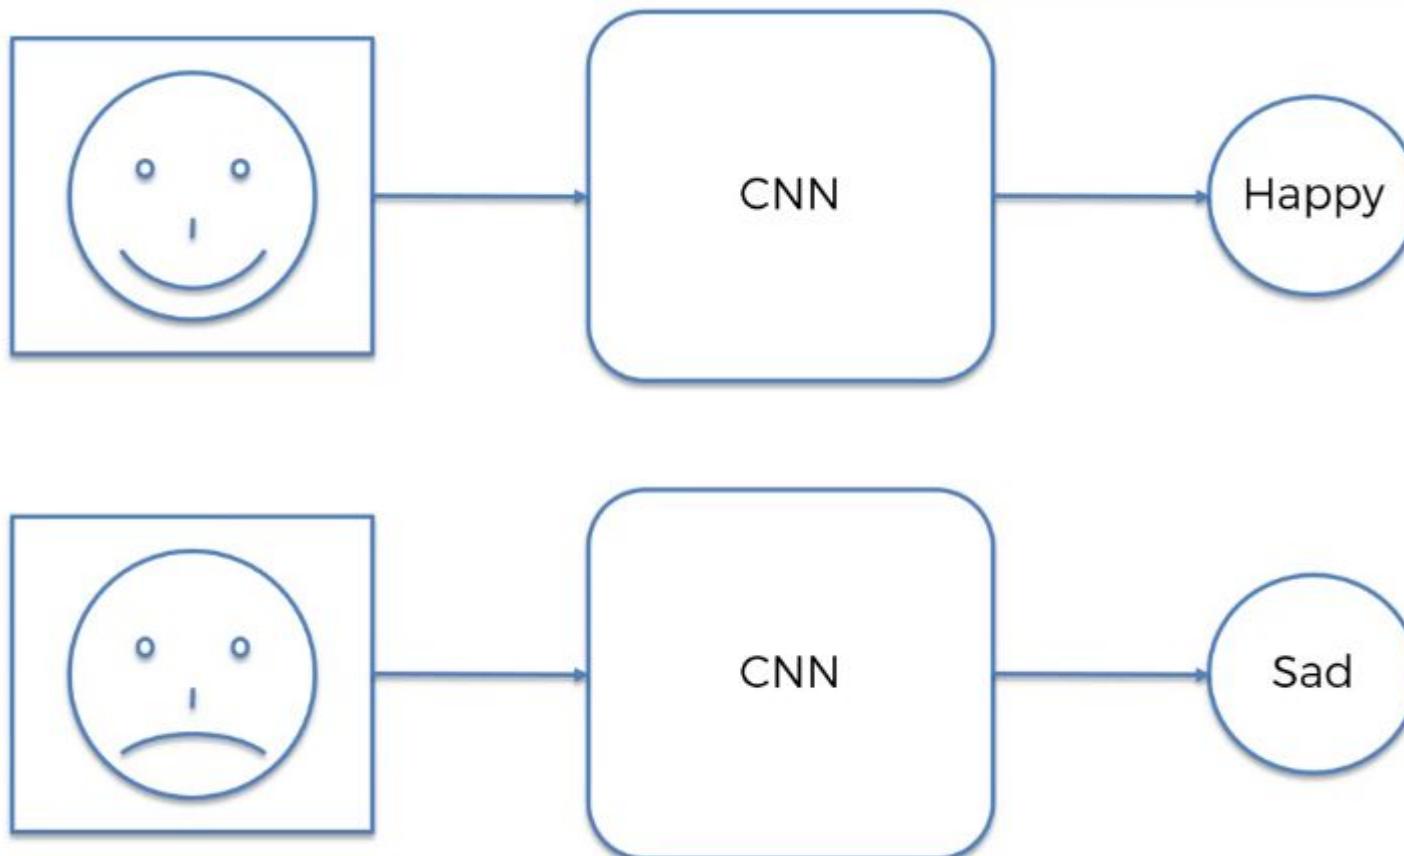
Google Facebook



CNNs



CNNs



CNNs

B / W Image 2x2px

Pixel 1	Pixel 2
Pixel 3	Pixel 4

2d array

Pixel 1 <small>0 ≤ pixel value ≤ 255</small>	Pixel 2 <small>0 ≤ pixel value ≤ 255</small>
Pixel 3 <small>0 ≤ pixel value ≤ 255</small>	Pixel 4 <small>0 ≤ pixel value ≤ 255</small>

Colored Image 2x2px

Pixel 1	Pixel 2
Pixel 3	Pixel 4

3d array

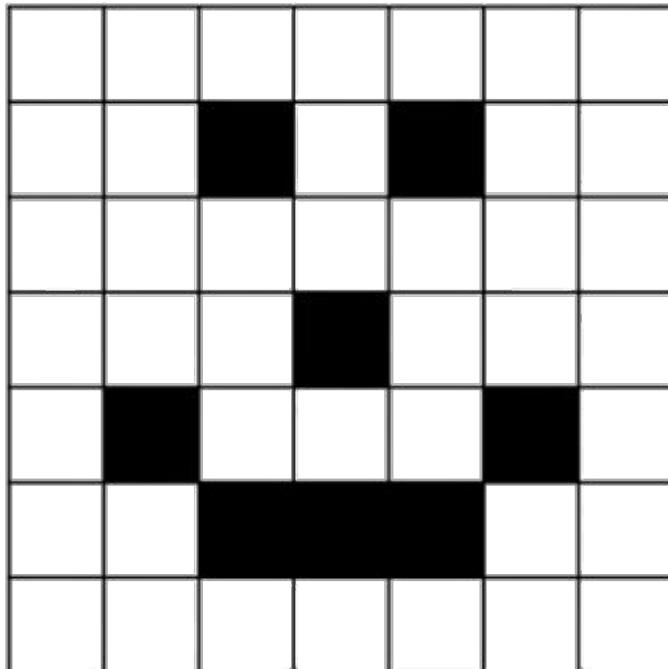
Red channel

Green channel

Blue channel

Pixel 1 <small>0 ≤ pixel value ≤ 255</small>	Pixel 2 <small>0 ≤ pixel value ≤ 255</small>
Pixel 3 <small>0 ≤ pixel value ≤ 255</small>	Pixel 4 <small>0 ≤ pixel value ≤ 255</small>

CNNs



0	0	0	0	0	0	0	0
0	1	0	0	0	1	0	0
0	0	0	0	0	0	0	0
0	0	0	1	0	0	0	0
0	1	0	0	0	1	0	0
0	0	1	1	1	0	0	0
0	0	0	0	0	0	0	0

Steps of a CNN

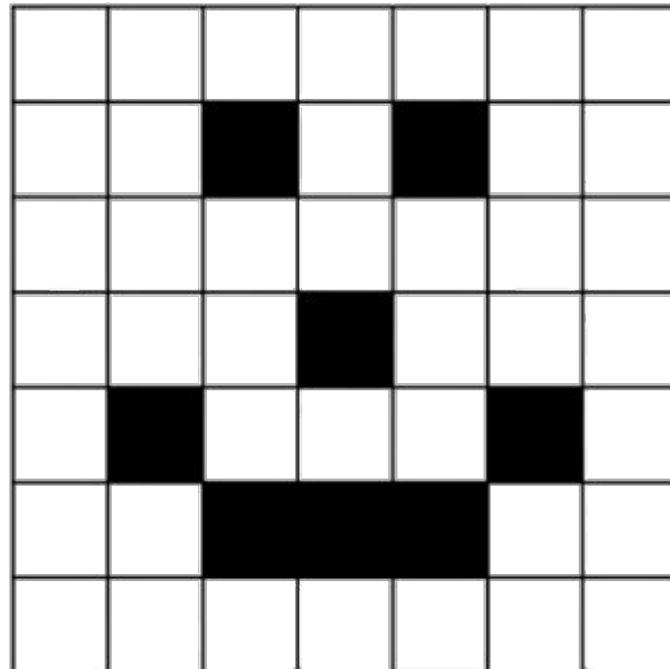
1. Convolution
2. Max Pooling
3. Flattening
4. Full Connection

Step 1 Convolution

CNN Step 1: Convolution

$$(f * g)(t) \stackrel{\text{def}}{=} \int_{-\infty}^{\infty} f(\tau) g(t - \tau) d\tau$$

CNN Step 1: Convolution



0	0	0	0	0	0	0	0
0	1	0	0	0	1	0	0
0	0	0	0	0	0	0	0
0	0	0	1	0	0	0	0
0	1	0	0	0	1	0	0
0	0	1	1	1	0	0	0
0	0	0	0	0	0	0	0

CNN Step 1: Convolution

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0

Input Image

0	0	1
1	0	0
0	1	1

Feature Detector

CNN Step 1: Convolution

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0

Input Image



0	0	1
1	0	0
0	1	1

Feature Detector

=

0			

Feature Map

CNN Step 1: Convolution

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0



Input Image

0	0	1
1	0	0
0	1	1

Feature Detector



0	1			

Feature Map

CNN Step 1: Convolution

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0



Input Image

0	0	1
1	0	0
0	1	1

Feature Detector



0	1	0		

Feature Map

CNN Step 1: Convolution

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0



Input Image

0	0	1
1	0	0
0	1	1

Feature Detector



0	1	0	0

Feature Map

Convolution

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0

Input Image



0	0	1
1	0	0
0	1	1

Feature Detector

=

0	1	0	0	0

Feature Map

CNN Step 1: Convolution

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0



Input Image

0	0	1
1	0	0
0	1	1

Feature Detector



0	1	0	0	0
0				

Feature Map

CNN Step 1: Convolution

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0



Input Image

0	0	1
1	0	0
0	1	1

Feature Detector



0	1	0	0	0
0	1			

Feature Map

CNN Step 1: Convolution

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0



Input Image

0	0	1
1	0	0
0	1	1

Feature Detector



0	1	0	0	0
0	1	1		

Feature Map

CNN Step 1: Convolution

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0



Input Image

0	0	1
1	0	0
0	1	1

Feature Detector



0	1	0	0	0
0	1	1	1	

Feature Map

CNN Step 1: Convolution

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0



Input Image

0	0	1
1	0	0
0	1	1

Feature Detector



0	1	0	0	0
0	1	1	1	0

Feature Map

CNN Step 1: Convolution

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0



Input Image

0	0	1
1	0	0
0	1	1

Feature Detector



0	1	0	0	0
0	1	1	1	0
1				

Feature Map

CNN Step 1: Convolution

0	0	0	0	0	0	0	0
0	1	0	0	0	1	0	0
0	0	0	0	0	0	0	0
0	0	0	1	0	0	0	0
0	1	0	0	0	1	0	0
0	0	1	1	1	0	0	0
0	0	0	0	0	0	0	0



Input Image

0	0	1
1	0	0
0	1	1

Feature Detector



0	1	0	0	0
0	1	1	1	0
1	0			

Feature Map

CNN Step 1: Convolution

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0



Input Image

0	0	1
1	0	0
0	1	1

Feature Detector



0	1	0	0	0
0	1	1	1	0
1	0	1		

Feature Map

CNN Step 1: Convolution

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0



Input Image

0	0	1
1	0	0
0	1	1

Feature Detector

=

0	1	0	0	0
0	1	1	1	0
1	0	1	2	

Feature Map

CNN Step 1: Convolution

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0



Input Image

0	0	1
1	0	0
0	1	1

Feature Detector



0	1	0	0	0
0	1	1	1	0
1	0	1	2	1

Feature Map

Convolution

0	0	0	0	0	0	0	0
0	1	0	0	0	1	0	0
0	0	0	0	0	0	0	0
0	0	0	1	0	0	0	0
0	1	0	0	0	1	0	0
0	0	1	1	1	0	0	0
0	0	0	0	0	0	0	0



Input Image

0	0	1
1	0	0
0	1	1

Feature Detector



0	1	0	0	0
0	1	1	1	0
1	0	1	2	1
1				

Feature Map

CNN Step 1: Convolution

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0



Input Image

0	0	1
1	0	0
0	1	1

Feature Detector



0	1	0	0	0
0	1	1	1	0
1	0	1	2	1
1	4			

Feature Map

CNN Step 1: Convolution

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0



Input Image

0	0	1
1	0	0
0	1	1

Feature Detector



0	1	0	0	0
0	1	1	1	0
1	0	1	2	1
1	4	2		

Feature Map

CNN Step 1: Convolution

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0



Input Image

0	0	1
1	0	0
0	1	1

Feature Detector



0	1	0	0	0
0	1	1	1	0
1	0	1	2	1
1	4	2	1	

Feature Map

CNN Step 1: Convolution

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0



Input Image

0	0	1
1	0	0
0	1	1

Feature Detector



0	1	0	0	0
0	1	1	1	0
1	0	1	2	1
1	4	2	1	0

Feature Map

CNN Step 1: Convolution

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0



Input Image

0	0	1
1	0	0
0	1	1

Feature Detector



0	1	0	0	0
0	1	1	1	0
1	0	1	2	1
1	4	2	1	0
0				

Feature Map

CNN Step 1: Convolution

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0



Input Image

0	0	1
1	0	0
0	1	1

Feature Detector



0	1	0	0	0
0	1	1	1	0
1	0	1	2	1
1	4	2	1	0
0	0			

Feature Map

CNN Step 1: Convolution

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0



Input Image

0	0	1
1	0	0
0	1	1

Feature Detector



0	1	0	0	0
0	1	1	1	0
1	0	1	2	1
1	4	2	1	0
0	0	1		

Feature Map

CNN Step 1: Convolution

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0



Input Image

0	0	1
1	0	0
0	1	1

Feature Detector



0	1	0	0	0
0	1	1	1	0
1	0	1	2	1
1	4	2	1	0
0	0	1	2	

Feature Map

CNN Step 1: Convolution

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0



Input Image

0	0	1
1	0	0
0	1	1

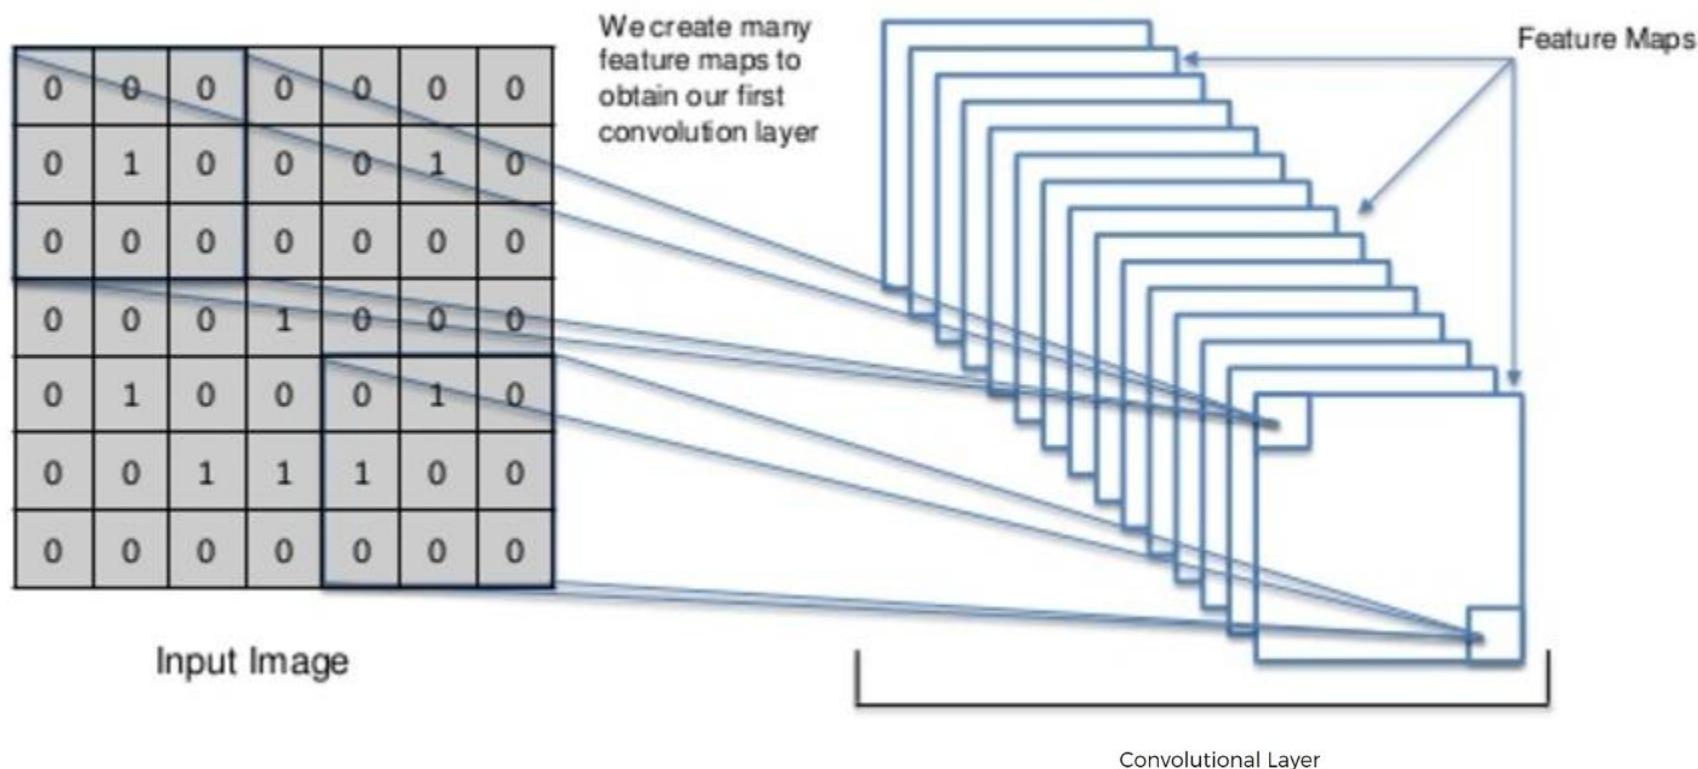
Feature Detector



0	1	0	0	0
0	1	1	1	0
1	0	1	2	1
1	4	2	1	0
0	0	1	2	1

Feature Map

CNN Step 1: Convolution

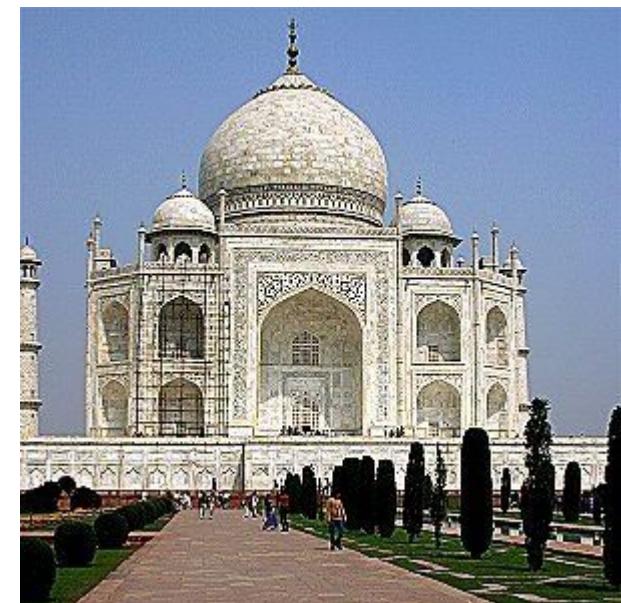


Convolution result examples



Sharpen

0	0	0	0	0
0	0	-1	0	0
0	-1	5	-1	0
0	0	-1	0	0
0	0	0	0	0



Convolution result examples



Blur

0	0	0	0	0
0	1	1	1	0
0	1	1	1	0
0	1	1	1	0
0	0	0	0	0

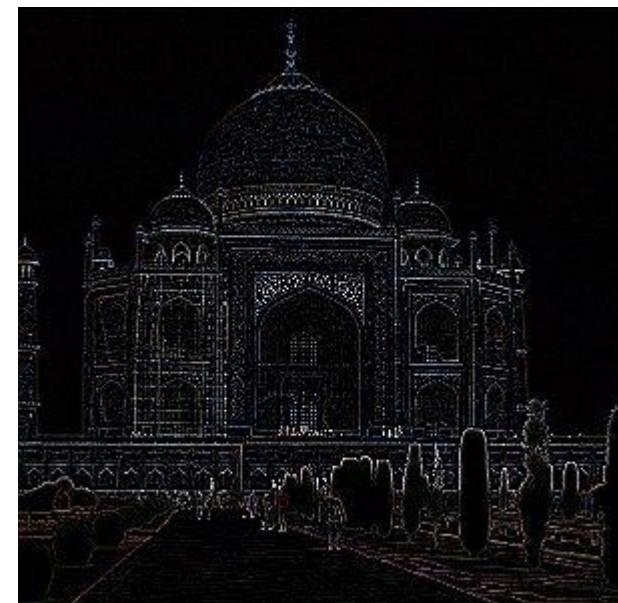


Convolution result examples



Edge Detect

0	1	0
1	-4	1
0	1	0

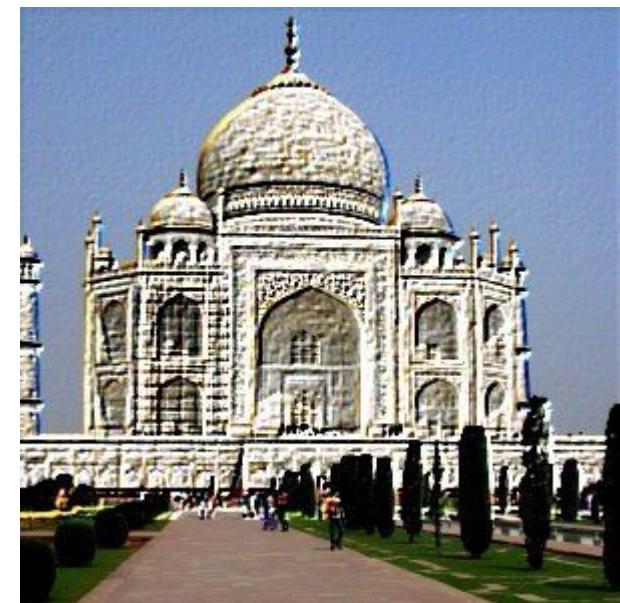


Convolution result examples



Emboss

-2	-1	0
-1	1	1
0	1	2



color images convolution

0	0	0	0	0	0	0	...
0	156	155	156	158	158	158	...
0	153	154	157	159	159	159	...
0	149	151	155	158	159	159	...
0	146	146	149	153	158	158	...
0	145	143	143	148	158	158	...
...

Input Channel #1 (Red)

0	0	0	0	0	0	0	...
0	167	166	167	169	169	169	...
0	164	165	168	170	170	170	...
0	160	162	166	169	170	170	...
0	156	156	159	163	168	168	...
0	155	153	153	158	168	168	...
...

Input Channel #2 (Green)

0	0	0	0	0	0	0	...
0	163	162	163	165	165	165	...
0	160	161	164	166	166	166	...
0	156	158	162	165	166	166	...
0	155	155	158	162	167	167	...
0	154	152	152	157	167	167	...
...

Input Channel #3 (Blue)

-1	-1	1
0	1	-1
0	1	1

Kernel Channel #1

1	0	0
1	-1	-1
1	0	-1

Kernel Channel #2

0	1	1
0	1	0
1	-1	1

Kernel Channel #3

↓
308

+

↓
-498

+

↓
164

↑
Bias = 1

+ 1 = -25

Output

-25				...
				...
				...
				...
...

CNN Step 1: Convolution

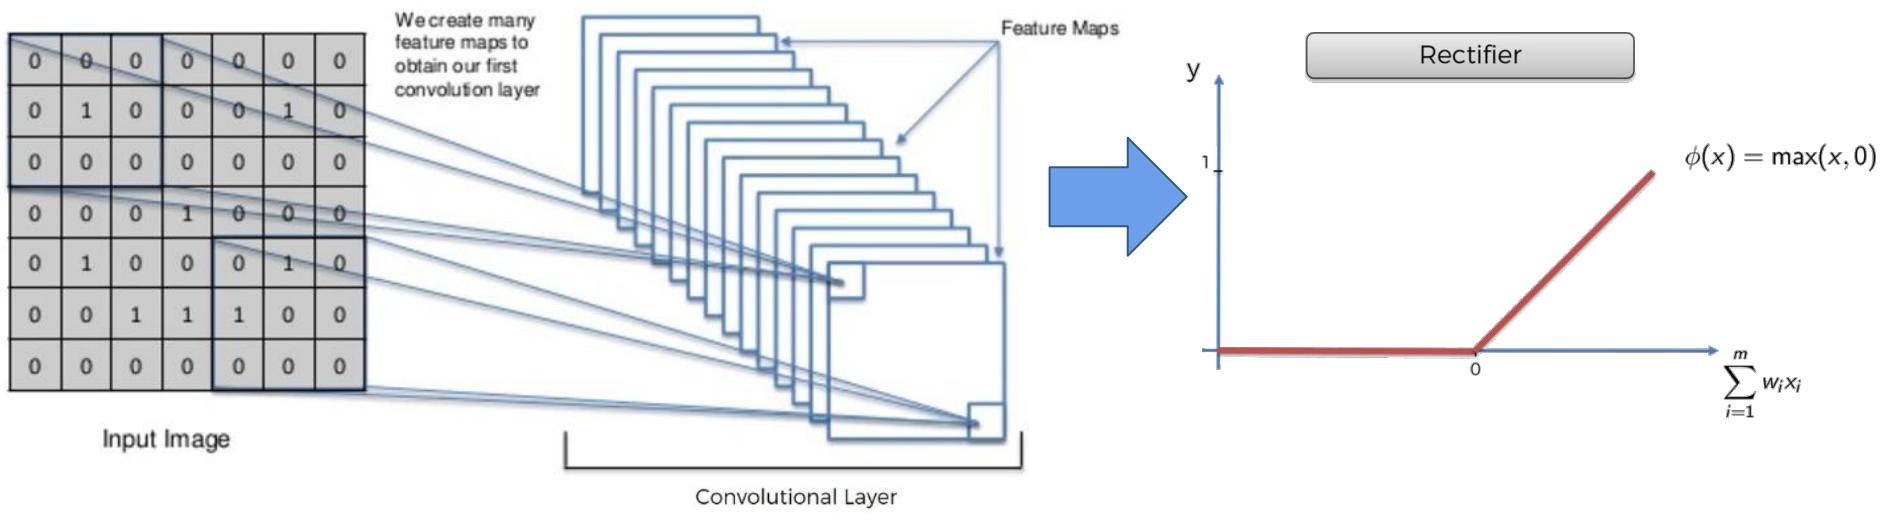
Searching for features in the images, using a feature detector, and putting them on a feature map that still preserves the spatial relationships between pixels.

Many times, we won't know what these characteristics mean, but they work.

Step 1(b)

ReLU Activation

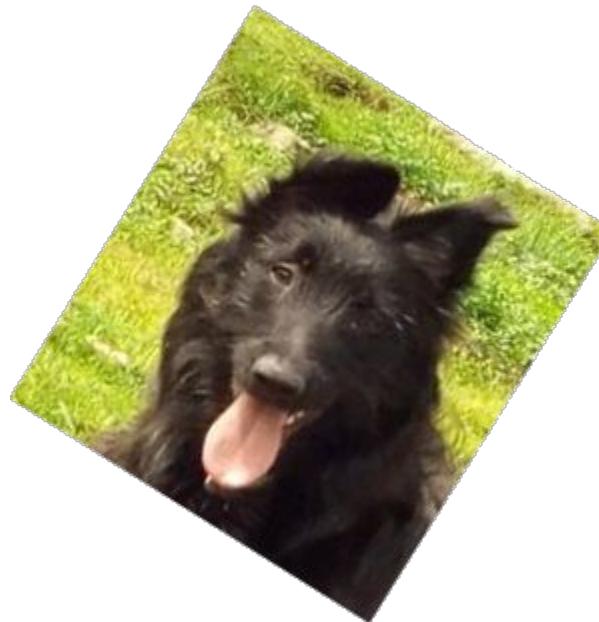
Step 1b: ReLU layer



We add NonLinearity to the output

Step 2 Max Pooling

Step 2: Pooling



Step 2: Pooling



Step 2: MAX Pooling

0	1	0	0	0
0	1	1	1	0
1	0	1	2	1
1	4	2	1	0
0	0	1	2	1

Feature Map

Step 2: MAX Pooling

0	1	0	0	0
0	1	1	1	0
1	0	1	2	1
1	4	2	1	0
0	0	1	2	1

Feature Map

Max Pooling

Pooled Feature Map

Step 2: MAX Pooling

0	1	0	0	0
0	1	1	1	0
1	0	1	2	1
1	4	2	1	0
0	0	1	2	1

Feature Map

Max Pooling



1		

Pooled Feature Map

Step 2: MAX Pooling

0	1	0	0	0
0	1	1	1	0
1	0	1	2	1
1	4	2	1	0
0	0	1	2	1

Feature Map

Max Pooling



1	1	

Pooled Feature Map

Step 2: MAX Pooling

0	1	0	0	0
0	1	1	1	0
1	0	1	2	1
1	4	2	1	0
0	0	1	2	1

Feature Map

Max Pooling

1	1	0

Pooled Feature Map

Step 2: MAX Pooling

0	1	0	0	0
0	1	1	1	0
1	0	1	2	1
1	4	2	1	0
0	0	1	2	1

Feature Map

Max Pooling



1	1	0
4		

Pooled Feature Map

Step 2: MAX Pooling

0	1	0	0	0
0	1	1	1	0
1	0	1	2	1
1	4	2	1	0
0	0	1	2	1

Feature Map

Max Pooling



1	1	0
4	2	

Pooled Feature Map

Step 2: MAX Pooling

0	1	0	0	0
0	1	1	1	0
1	0	1	2	1
1	4	2	1	0
0	0	1	2	1

Feature Map

Max Pooling

1	1	0
4	2	1

Pooled Feature Map

Step 2: MAX Pooling

0	1	0	0	0
0	1	1	1	0
1	0	1	2	1
1	4	2	1	0
0	0	1	2	1

Feature Map

Max Pooling



1	1	0
4	2	1
0		

Pooled Feature Map

MAX Pooling

0	1	0	0	0
0	1	1	1	0
1	0	1	2	1
1	4	2	1	0
0	0	1	2	1

Feature Map

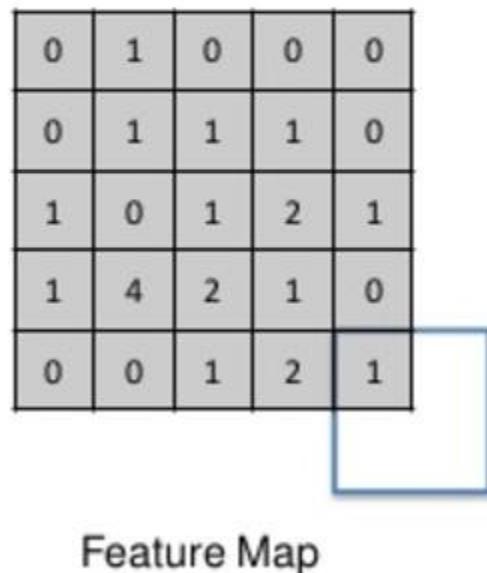
Max Pooling



1	1	0
4	2	1
0	2	

Pooled Feature Map

Step 2: MAX Pooling



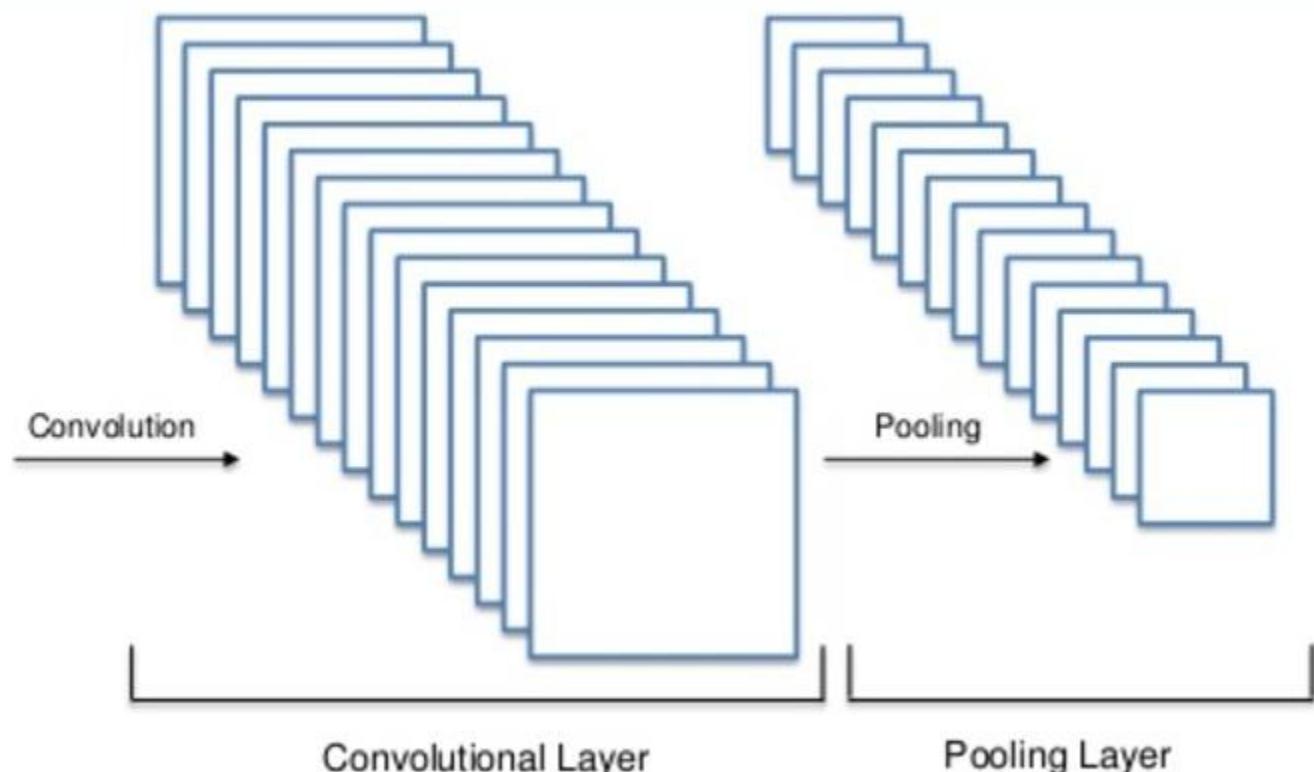
1	1	0
4	2	1
0	2	1

Pooled Feature Map

Step 1 and 2: Conv + pooling

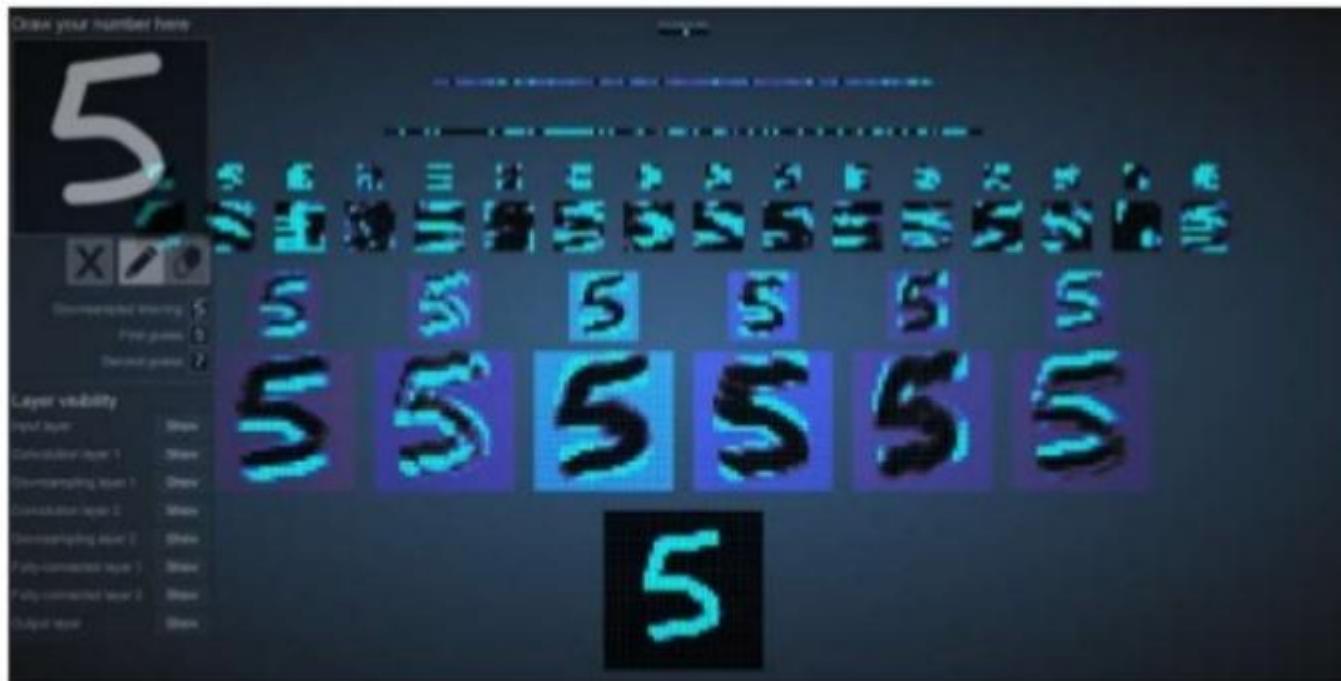
0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0

Input Image



Step 1 and 2: Conv + pooling

- Interesting paper about poolings:
http://www.ais.uni-bonn.de/papers/icann2010_maxpool.pdf
- Convolution and pooling live! example:
<http://scs.ryerson.ca/~aharley/vis/conv/flat.html>
-

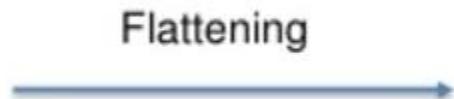


Step 3 Flattening

Step 3: Flattening

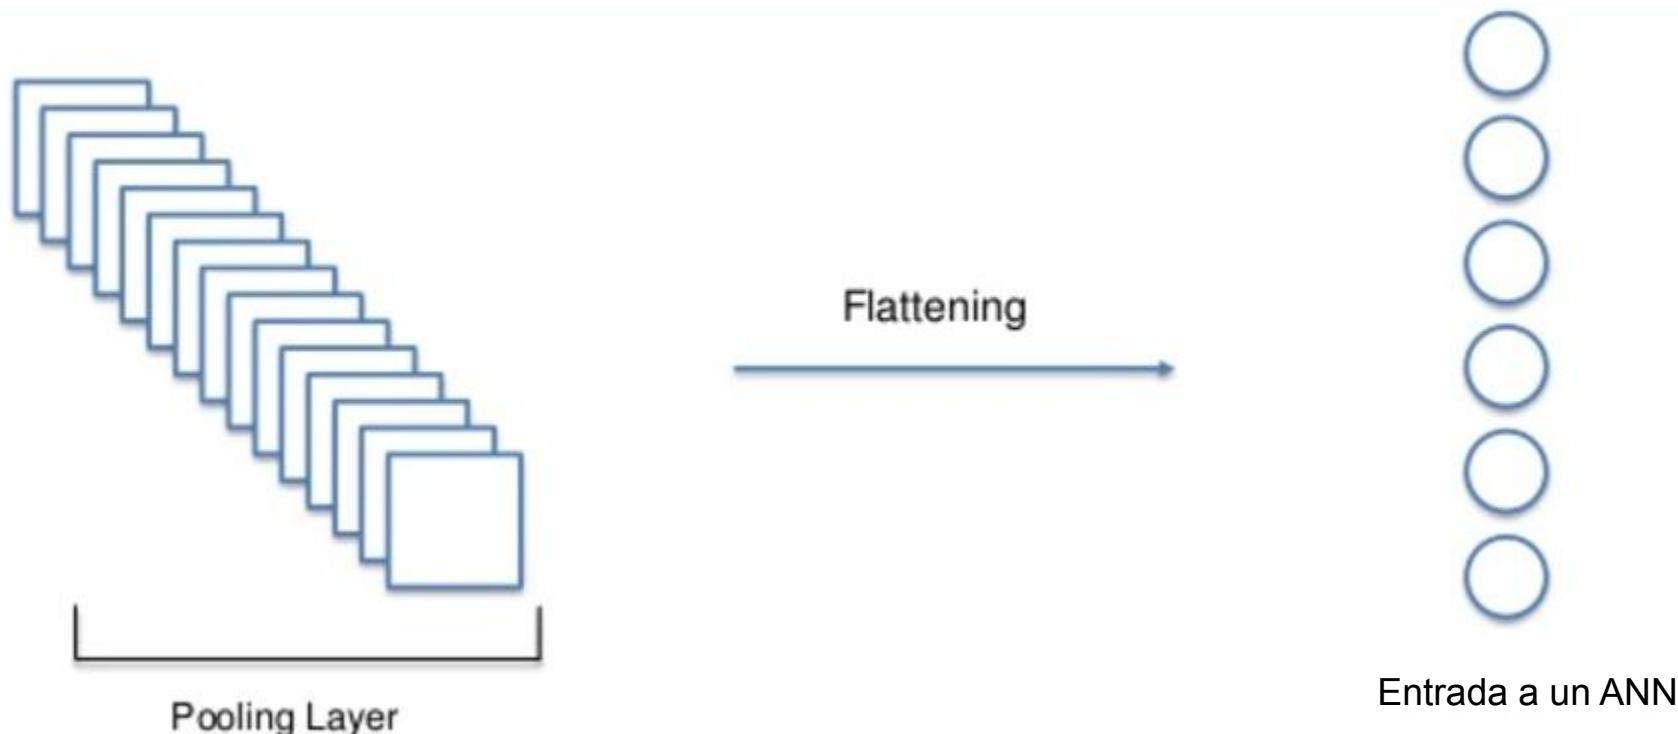
1	1	0
4	2	1
0	2	1

Pooled Feature Map



1
1
0
4
2
1
0
2
1

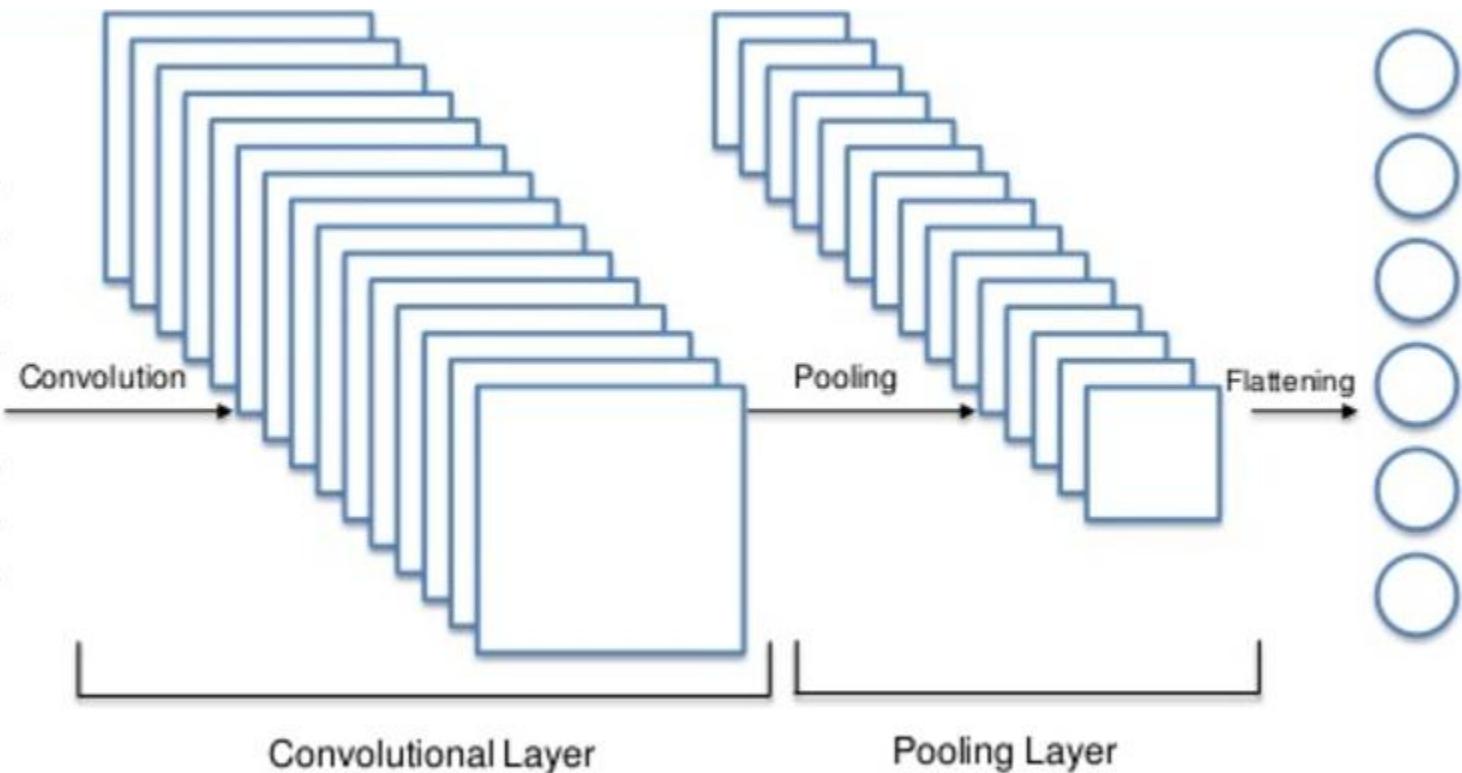
Step 3: Flattening



steps 1+2+3: conv.+pool.+flat.

0	0	0	0	0	0	0	0
0	1	0	0	0	1	0	
0	0	0	0	0	0	0	
0	0	0	1	0	0	0	
0	1	0	0	0	1	0	
0	0	1	1	1	0	0	
0	0	0	0	0	0	0	

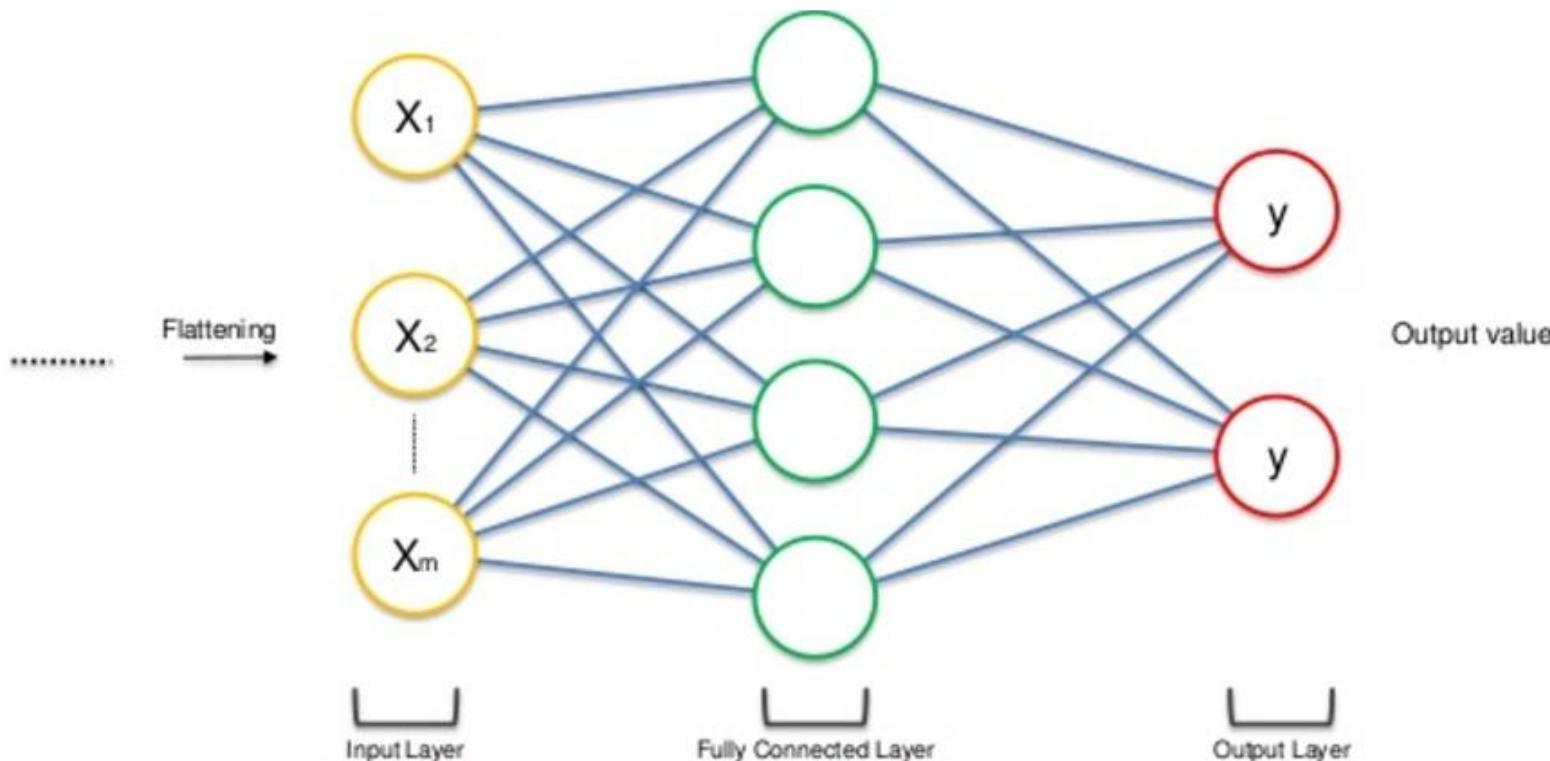
Input Image



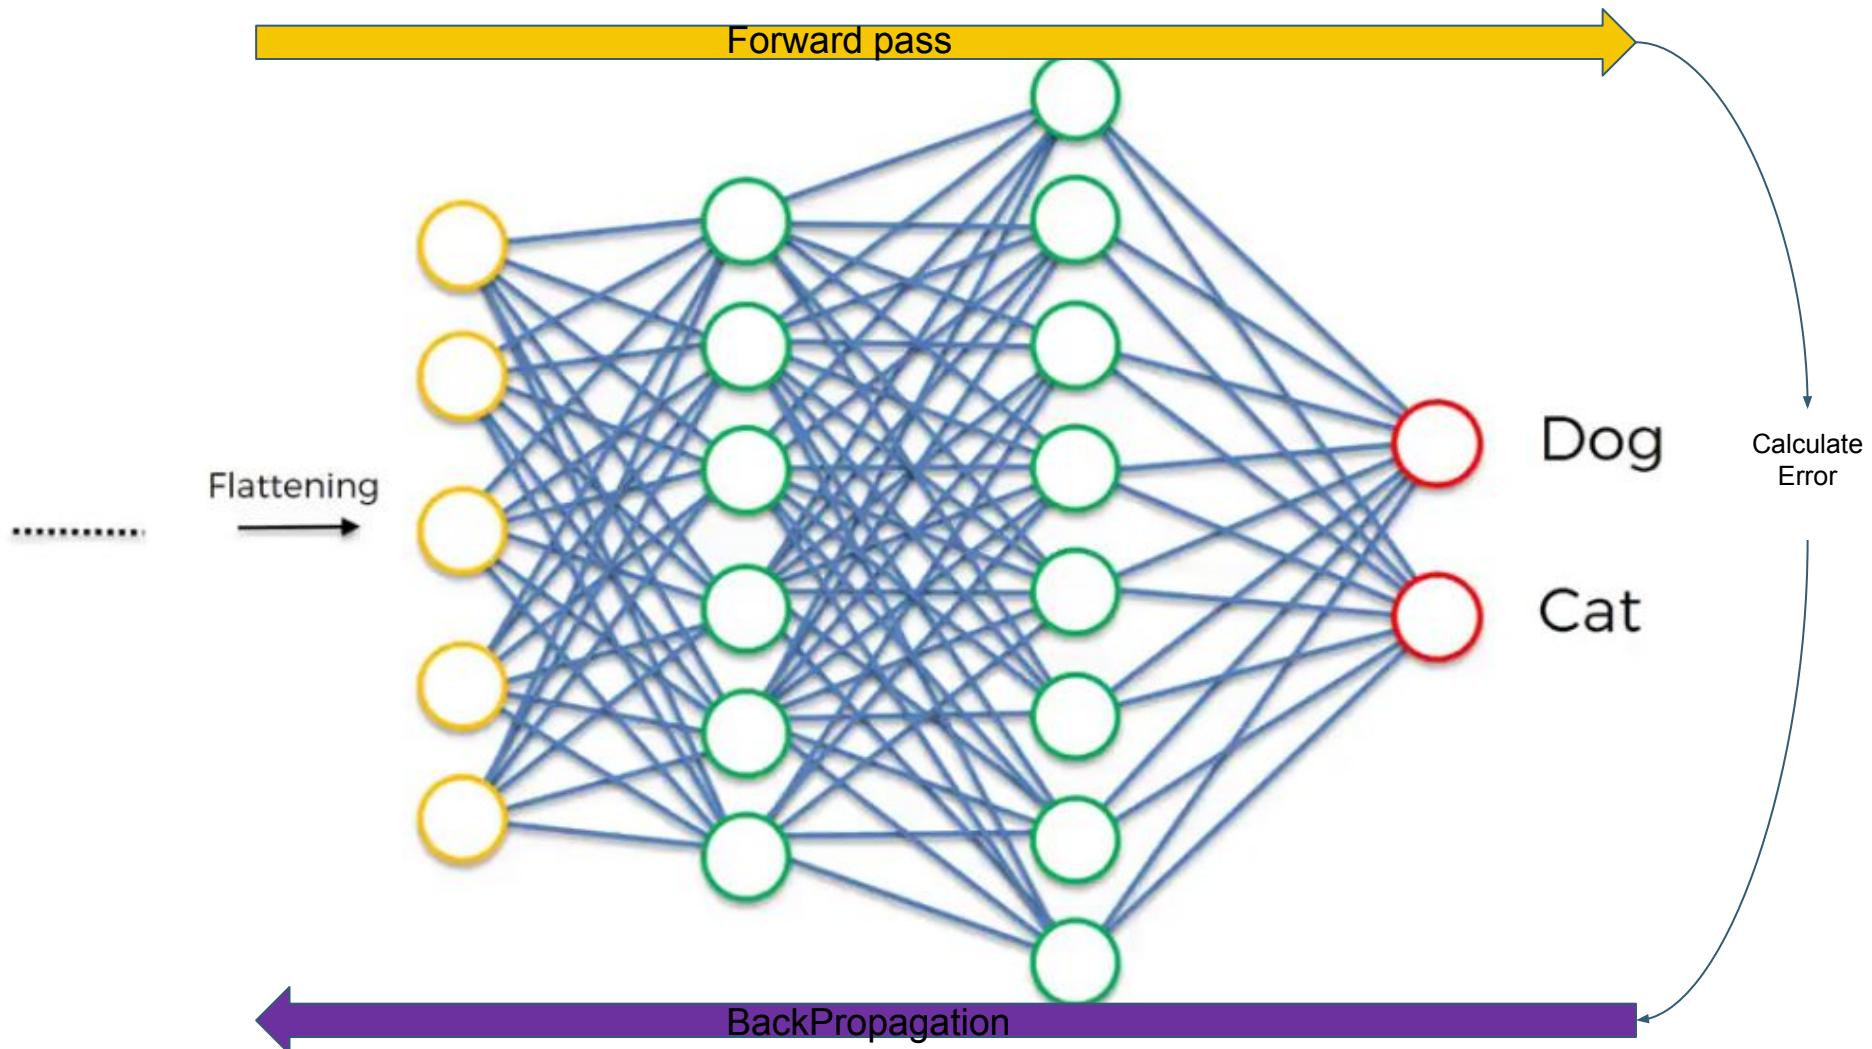
Step 4

Full Connection

Step 4: Fully connected ANN

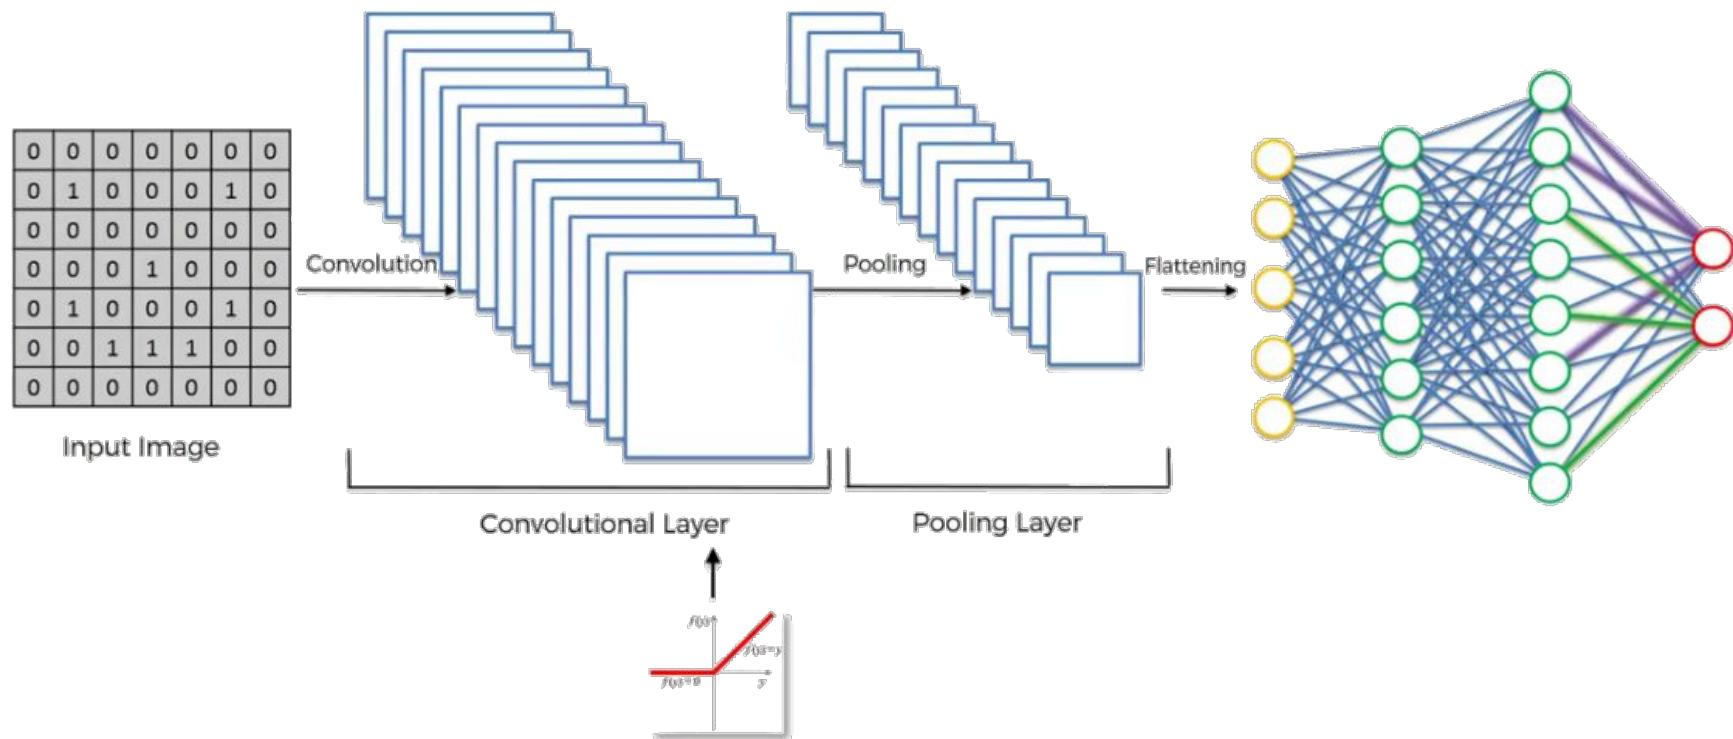


Step 4: Fully connected ANN



Summary

Summary



CNN Tensorflow utils

Image preprocessing with TF

- Preprocessing steps:
 - Read the picture files —
 - decode JPEG content to RGB pixels —
 - Convert these into floating-point tensors —
 - Rescale the pixels to be between [0,1] —
- Image Agumentation
 - rotate —
 - width/height shift —
 - shear —
 - zoom —
 - horizontal/vertical flip —
- Luckily, Keras has a `ImageDataGenerator` that makes —
- This generator can also —

Image preprocessing with TF

```

train_datagen = ImageDataGenerator(
    rescale=1./255,
    rotation_range=40,
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True, )

test_datagen = ImageDataGenerator(rescale=1./255)

train_generator = train_datagen.flow_from_directory(
    train_dir,
    target_size=(150, 150),
    batch_size=32,
    class_mode='binary')

validation_generator = test_datagen.flow_from_directory(
    validation_dir,
    target_size=(150, 150),
    batch_size=32,
    class_mode='binary')

```

Target
directory

train_dir,

target_size=(150, 150),

batch_size=32,

class_mode='binary')

Resizes all images to 150 × 150

Note that the validation data shouldn't be augmented!

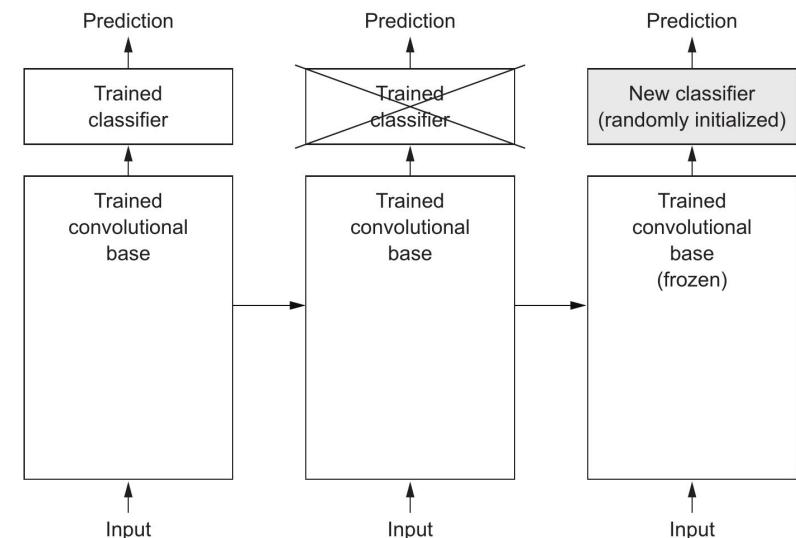
Because you use binary_crossentropy loss, you need binary labels.



Using a pretrained convnet

- It is common to use a pretrained network to extract characteristics from images
 - A pretrained network is a network that was trained previously on a huge dataset
- Usually, the features learned by a pretrained network are general enough to use them in many different vision problems
- Usually, we change the “fully connected” layer and “freeze” the convolutional part
- There are several pretrained networks
 - VGG16
 - MobileNet
 - ResNet
 - ...

```
from keras.applications import VGG16
conv_base = VGG16(weights='imagenet',
                  include_top=False,
                  input_shape=(150, 150, 3))
```



Exercises

1. Crear y entrenar una red neuronal que diferencie perros de gatos con Keras y CNN

1.1. Aprenderemos a crear todo lo explicado con las CNN

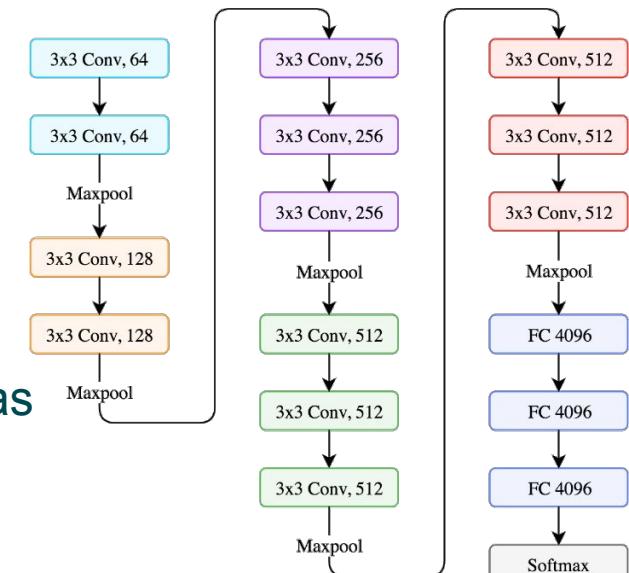
- 1.1.1. Capa Convolucional + ReLU
- 1.1.2. Max Pooling
- 1.1.3. Flattening
- 1.1.4. Red fully connected

2. Utilizar redes neuronales preentrenadas (VGG16)

2.1. Aprenderemos a utilizar redes ya creadas y entrenadas

- 2.1.1. VGG16

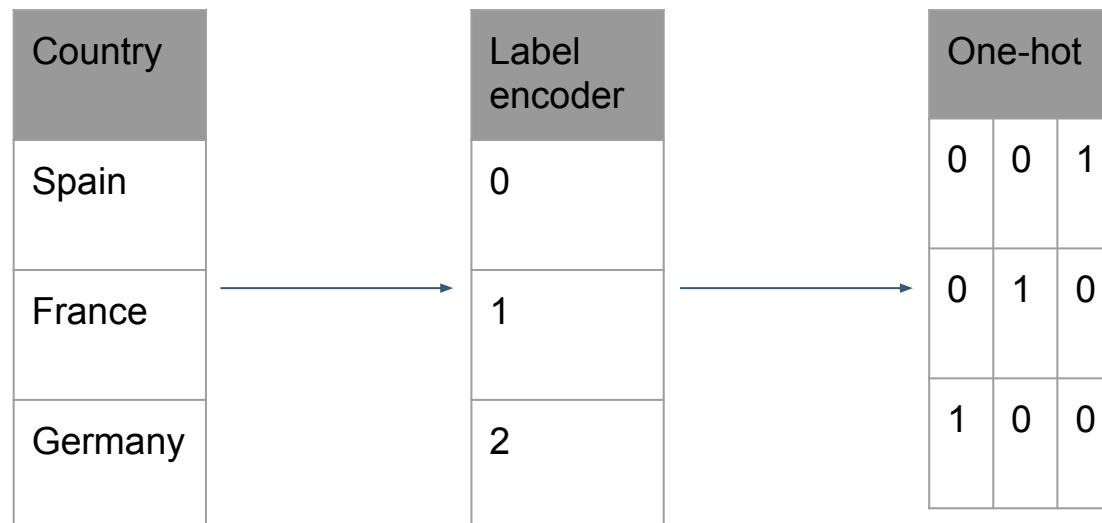
3. perros y gatos con resnet



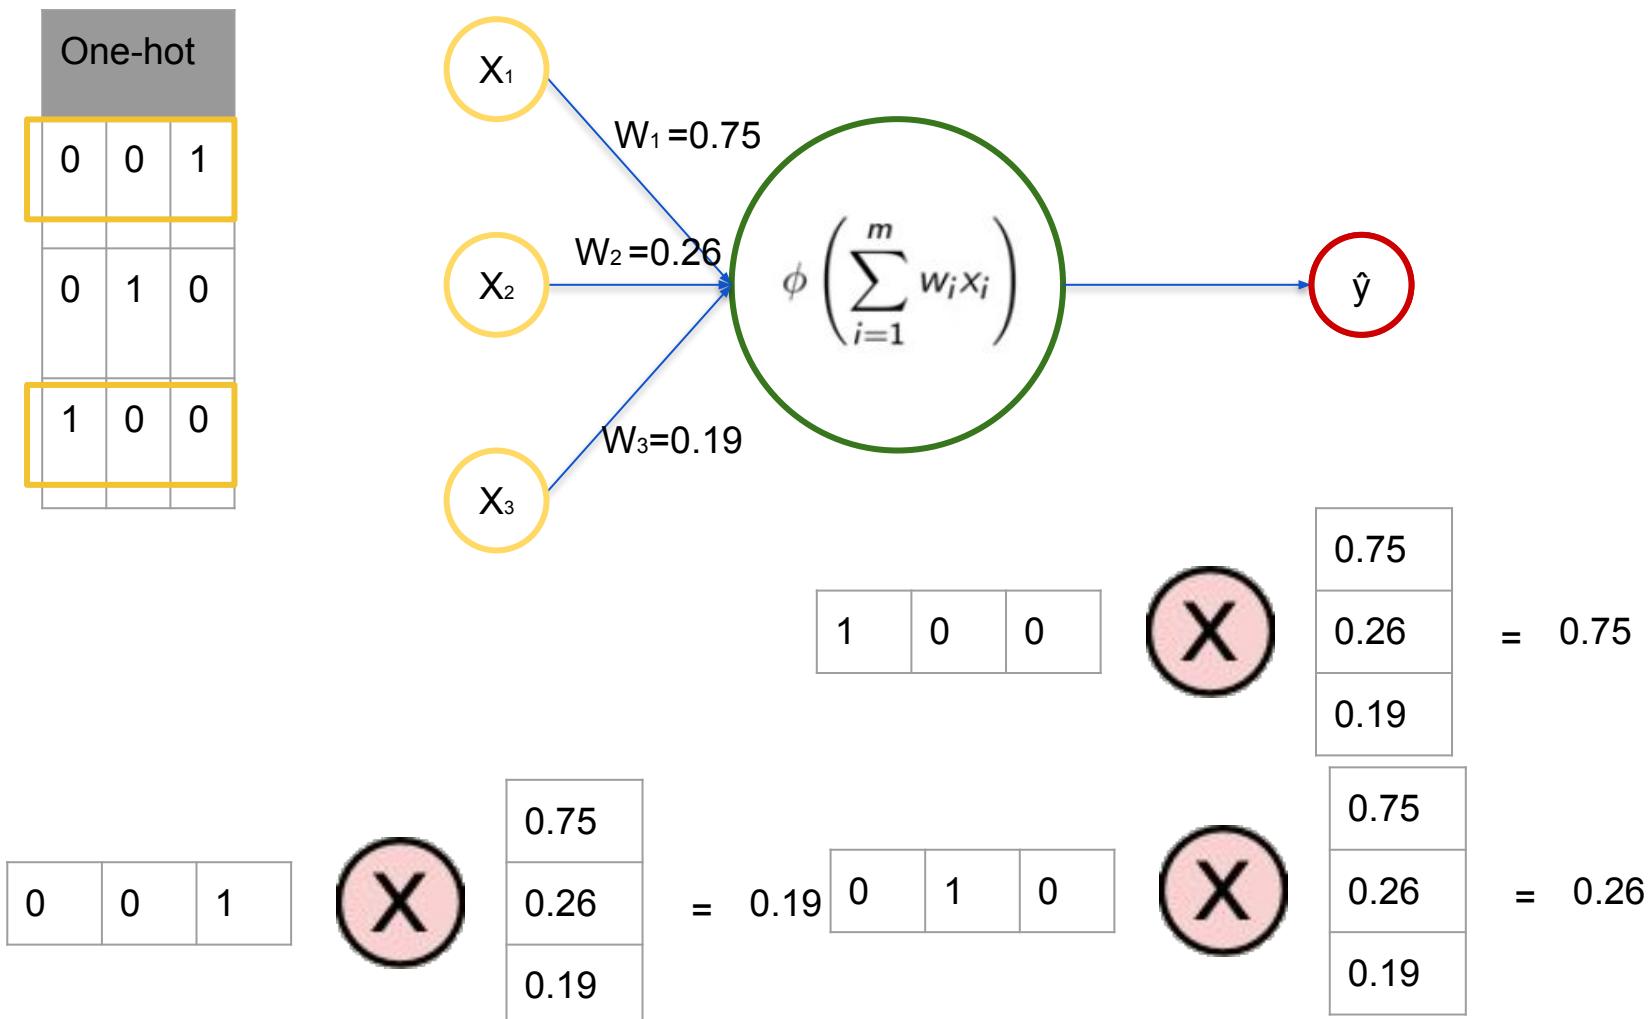
Apendix 1

Embedding Layer

- Se utiliza cuando hay datos categóricos

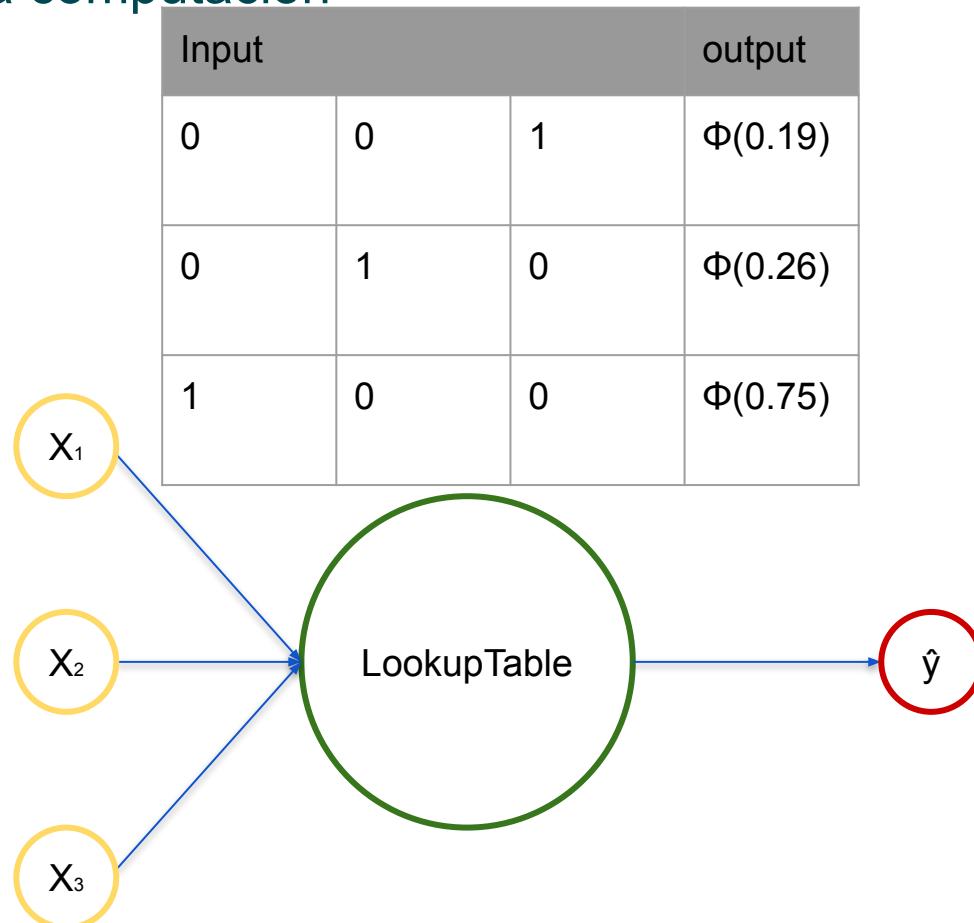


Embedding Layer



Embedding Layer

- No es más que un “lookup table” para datos categóricos
- así ahorra computación



Eskerrik asko
Muchas gracias
Thank you

Ekhi Zugasti
ezugasti@mondragon.edu

Loramendi, 4. Apartado 23
20500 Arrasate – Mondragon
T. 943 71 21 85
info@mondragon.edu